

xlog: A GPU-Native Logic Programming Language for Unified Symbolic Reasoning

The xlog Project
v0.9.2 Technical Whitepaper

June 2026

Abstract

We present xlog, a GPU-native logic programming language that unifies five reasoning paradigms—deterministic Datalog, probabilistic inference, epistemic reasoning over world views, SAT/MaxSAT solving, and differentiable neural-symbolic learning—on a single CUDA runtime, keeping all semantic state device-resident with zero host-device transfers on production paths. Its central contribution is a fully GPU-native knowledge-compilation pipeline (provenance $\rightarrow$ CNF $\rightarrow$ Decision-DNNF $\rightarrow$ arithmetic circuit) in which every compiled circuit is verified, by an on-GPU CDCL solver, to be logically equivalent to its source formula—turning circuit correctness into a machine-checked certificate rather than an assumption. xlog is implemented in Rust with 23 CUDA kernel files and offers zero-copy interoperability with PyTorch, JAX, and Arrow consumers via DLPack and the Arrow C Device interface. Its v0.9.x line adds epistemic logic programming—**know/possible** reasoning over world views under founded (FAEEL) and Gelfond-style semantics, including finite nested modal chains, epistemic integrity constraints, and recursive epistemic execution—executed GPU-natively through a Generate-Propagate-Test plan that reuses the same solver and join machinery. We report ablations on a single development GPU: circuit caching yields a $2.74\times$ end-to-end training speedup (95% CI [2.29, 3.18]) on MNIST addition by amortizing compilation across iterations, and a cost-aware worst-case-optimal join subsystem yields a $27.96\times$ geometric-mean speedup over xlog’s own binary-join baseline on skewed program-analysis workloads. These are ablations against xlog’s own baselines; we make no head-to-head comparison against external systems, and the contribution is the unified, transfer-free architecture rather than a claim of being fastest.

Contents

1	Introduction	3
2	Architecture	5
2.1	System Overview	5
2.2	Compilation Pipeline	8
2.3	GPU Residency Model	8
2.4	IR Stack	9
3	The xlog Language	10
3.1	Design Philosophy	10
3.2	Type System	11
3.3	Expressions and Arithmetic	11

3.4	User-Defined Functions	12
3.5	Modules and Imports	13
3.6	Aggregations and Constraints	14
3.7	Language Completeness over Finite Domains	14
4	GPU-Native Datalog Evaluation	15
4.1	Semi-Naive Evaluation on GPU	16
4.2	Kernel Design	16
4.3	Adaptive Join Planning	17
4.4	Worst-Case Optimal Joins	18
4.5	Runtime Optimization Layer	18
4.6	Reversible Symbols	18
5	Probabilistic Inference on GPU	19
5.1	Approach	19
5.2	GPU Knowledge Compilation Pipeline	19
5.3	Circuit Caching	20
5.4	Monte Carlo Sampling	21
5.5	Probabilistic Aggregates	21
5.6	Well-Founded Semantics	22
6	Neural-Symbolic Bridge	22
6.1	Neural Predicates	22
6.2	End-to-End Training Loop	24
6.3	Term Embeddings (v0.5.0)	26
6.4	Differentiable ILP (Beta)	27
7	Epistemic Reasoning	28
7.1	Modal Surface and the Epistemic IR	28
7.2	World-View Semantics	28
7.3	GPU-Native Execution	29
7.4	Solver and Probabilistic Integration	29
8	Interoperability	29
8.1	DLPack: Zero-Copy GPU Tensor Sharing	29
8.2	Arrow IPC: Columnar Export with Dictionary Encoding	30
8.3	Python Bindings	30
8.4	PyTorch Autograd Integration	30

9	Evaluation	30
9.1	Methodology	31
9.2	Absolute Performance	31
9.2.1	Deterministic Subsystem	31
9.2.2	Worst-Case Optimal Joins	31
9.2.3	Runtime Optimization Layer	32
9.2.4	Probabilistic Subsystem	32
9.2.5	Neural-Symbolic Subsystem	32
9.3	Qualitative Comparison	33
9.4	CUDA Certification	34
10	Related Work	35
10.1	Logic Programming Languages	35
10.2	GPU Datalog	35
10.3	Probabilistic Logic Programming	36
10.4	GPU SAT	36
10.5	Differentiable ILP	36
10.6	Positioning	37
11	Limitations and Future Work	37
11.1	Current Limitations	37
11.2	Future Work	38
A	Performance Contract	39

1 Introduction

Symbolic AI and neural AI have followed divergent engineering paths. Datalog engines, probabilistic logic systems such as ProbLog, and inductive logic programming (ILP) frameworks are implemented as CPU-bound interpreters or compilers, processing relations and proofs in main memory. Deep learning frameworks—PyTorch, JAX—execute dense tensor computations on GPUs via highly optimized CUDA kernels. When researchers combine the two paradigms, as in DeepProbLog [1] or NeurASP [2], the symbolic component remains on the CPU while the neural component runs on the GPU. Every training iteration transfers data across the PCIe bus: the CPU-side logic engine materializes query results, ships them to the GPU for gradient computation, then pulls gradients back to update symbolic parameters. At scale—millions of ground atoms, thousands of training steps—these host-device transfers dominate wall-clock time and memory bandwidth, becoming the primary bottleneck rather than the inference or learning computation itself.

The gap is architectural. Existing systems address individual reasoning tasks on the GPU in isolation: GPU-accelerated Datalog evaluation [3, 4], GPU SAT solvers [5], or differentiable logic on CPU with GPU-side neural networks [6]. No single system performs deterministic Datalog evaluation,

probabilistic inference via knowledge compilation, SAT/MaxSAT verification, and differentiable neural-symbolic training entirely on the GPU with zero host-device data transfers in production paths. The absence of such a platform forces practitioners into multi-system pipelines—a Datalog engine for rule evaluation, a separate probabilistic reasoner, a Python training loop bridging CPU logic to GPU tensors—each with its own data format, memory model, and failure modes.

xlog addresses this gap as a GPU-native logic programming language whose compiler and runtime span five reasoning paradigms: deterministic logic (semi-naive evaluation with stratified negation), probabilistic inference (exact knowledge compilation and Monte Carlo sampling), epistemic reasoning (world-view semantics over multiple stable models), SAT/MaxSAT solving (GPU CDCL with proof certificates), and neural-symbolic learning (differentiable training with PyTorch interoperability). The system is implemented in Rust with 23 custom CUDA kernel files (17.7K lines of device code) organized into a layered crate architecture. The compilation pipeline transforms Datalog source into a relational intermediate representation (RIR), lowers probabilistic programs through a propositional intermediate representation (PIR) into CNF, compiles decision-DNNF circuits via D4 [7], and encodes the result in a GPU-resident circuit format (XGCF) for forward and backward evaluation. All semantic data structures—fact stores, circuit nodes, solver state, gradient buffers—remain GPU-resident during execution. Host involvement is limited to orchestration, I/O, and compilation.

The principal contributions of this paper are:

- **A typed logic programming language for GPU execution.** The `xlog-logic` crate implements a statically-typed logic language with user-defined functions, modules and imports, stratified aggregation, integrity constraints, probabilistic facts, and neural predicate declarations—all compiling through a single PEG parser, type inferencer, and SCC-based stratifier into the relational IR that feeds every downstream backend (Section 3). The surface is completeness-oriented over finite domains: finite lists, safe meta-predicates, explicit negation-as-failure, magic-set query planning, and probabilistic aggregates all normalize into the same typed relational paths rather than a CPU term interpreter.
- **GPU-resident semi-naive Datalog evaluation.** The `xlog-runtime` and `xlog-cuda` crates implement relational algebra operators (hash join, radix sort, filter, deduplication, set difference, grouped aggregation) as CUDA kernels, executing fixed-point iteration entirely on the GPU with columnar storage and HISA indexing [3].
- **Worst-case-optimal multiway joins with cost-aware planning.** A hypergraph planner promotes eligible multiway rules to a GPU worst-case-optimal join (WCOJ) subsystem with cardinality/selectivity-aware variable ordering, recursive-SCC integration, and adaptive skew-triggered dispatch, with deterministic binary-join fallback preserved. On skewed paper-scale program-analysis workloads this yields a measured $27.96\times$ geometric-mean speedup over the binary-join baseline (Section 9). A runtime optimization layer adds persistent hash indexing, common-subexpression elimination, and adaptive re-optimization for repeated sessions.
- **GPU-native knowledge compilation with verified circuits.** The `xlog-prob` crate compiles probabilistic Datalog programs through a PIR-to-CNF-to-D4-to-XGCF pipeline, producing GPU-resident arithmetic circuits with compile-once/evaluate-many semantics; forward and backward passes run as level-parallel CUDA kernels for exact weighted model

counting and gradient computation without host transfers. Crucially, an on-GPU CDCL solver proves each compiled circuit logically equivalent to its source formula, so circuit correctness is a machine-checked certificate rather than an assumption. To our knowledge this is the first end-to-end GPU-resident knowledge-compilation pipeline with on-device equivalence verification, and it is the paper’s principal contribution.

- **GPU-native epistemic reasoning.** The v0.9.x line adds epistemic logic programming over world views: modal `know/possible` literals are preserved in an Epistemic IR and executed by a Generate–Propagate–Test plan whose phases run as CUDA kernels, with no CPU fallback on the accepted path. The accepted surface covers finite nested modal chains, epistemic integrity constraints, multi-predicate epistemic coupling, and recursive epistemic execution. It reuses the same on-GPU CDCL solver and worst-case-optimal join machinery as the deterministic and probabilistic backends (Section 7).
- **End-to-end differentiable neural-symbolic training.** The `xlog-neural` crate and `pyxlog` Python package connect compiled circuits to PyTorch’s autograd graph. Circuit structure depends on the logic program, not on neural network weights, so compiled XGCF templates are cached across training iterations. This circuit caching yields a measured $2.74\times$ end-to-end training speedup (95% CI: [2.29, 3.18]) by eliminating redundant D4 recompilation.
- **Zero-copy interoperability with ML frameworks.** The `xlog-cuda` crate exposes GPU-resident query results and gradient tensors via DLPack [8] capsules and Arrow [9] IPC/C Data interfaces, enabling direct consumption by PyTorch, cuDF, and other frameworks without data copies or device synchronization.
- **Differentiable ILP with GPU-resident credit assignment.** The `pyxlog dILP` trainer implements sparse GPU mask computation, a fully device-resident credit/loss path with zero host transfers, and a six-gate promotion pipeline (convergence, novel-rate audit, regression check, holdout F1 threshold, ambiguity scan, typed schema validation) for transactional rule induction. This subsystem is currently in beta.

The remainder of this paper is organized as follows. Section 2 presents the system architecture and crate decomposition. Section 3 describes the `xlog` language surface. Section 4 details GPU-native deterministic Datalog evaluation. Section 5 covers the probabilistic inference pipeline. Section 6 details the neural-symbolic bridge and differentiable training. Section 7 introduces epistemic reasoning over world views. Section 8 discusses interoperability with external frameworks. Section 9 presents evaluation results. Section 10 surveys related work. Section 11 discusses limitations and future directions.

2 Architecture

2.1 System Overview

`xlog`’s current workspace contains 15 Rust crates organized around a five-tier product/runtime hierarchy. Figure 1 summarizes the primary dependency spine and certification surface; two support crates discussed below (`xlog-induce` and `xlog-integration`) are omitted from the diagram for




Figure 1: xlog crate hierarchy. Tier 0 is the leaf foundation; each higher tier depends only on strictly lower tiers.

readability. The architecture still enforces strict layering: no crate depends on a crate in its own tier or a higher tier.

Tier 0 contains `xlog-core`, the leaf crate that defines shared scalar types (`ScalarType`, `Schema`, `AggOp`), the `KernelProvider` trait, error types, memory budgets, runtime configuration, and a bidirectional `SymbolTable` for dictionary-encoded strings. Every other crate in the workspace depends on `xlog-core`.

Tier 1 provides domain-specific intermediate representations and service providers that depend only on `xlog-core`. `xlog-ir` defines the relational IR node tree (`RirNode`), expression algebra, and `ExecutionPlan` structures. `xlog-cuda` wraps CUDA device management via `cudarc` [10], stages CUDA kernel modules (cubin-first with portable PTX fallback), implements the `CudaKernelProvider` with modular submodules (relational, filter, group-by, arithmetic, I/O, ILP, probabilistic, transfer, kernel loading), and exposes Arrow IPC/C Data and DLPack interoperability. `xlog-stats` provides `StatsManager` for compiler feedback and runtime tracking. `xlog-neural` defines the `NetworkRegistry`, `NetworkHandle`, `TensorSourceRegistry`, `NeuralBridge`, and `BatchCollector` types for neural-symbolic integration, with an optional PyO3 [11] feature gate for Python interop.

Tier 2 implements the three core subsystems. `xlog-logic` implements the xlog language frontend (see Section 3): PEG parser (`pest`), SCC-based stratifier, AST-to-RIR lowerer, and cost-aware optimizer. An internal rule-rewriting pass normalizes repeated patterns before stratification, invisible to user programs. xlog programs may also carry `#pragma` directives (e.g., `prob_engine`, `prob_cache`, `max_recursion_depth`) that influence compiler behavior; pragmas parse into dedicated AST nodes consumed by the compiler before lowering. Its production dependencies are `xlog-core`, `xlog-ir`, and `xlog-stats`. `xlog-runtime` provides the `Executor` (modular across delta computation, expression evaluation, join caching, node dispatch, recursive evaluation, and rewriting), versioned `RelationStore`, profiling, and query statistics. It depends on `xlog-core`, `xlog-ir`, `xlog-cuda`, and `xlog-stats`. `xlog-solve` implements GPU CDCL verification (complete SAT/UNSAT with on-GPU model and proof validation) and continuous local search for SAT/MaxSAT. It depends on `xlog-core` and `xlog-cuda`.

Tier 3 composes Tier 2 subsystems into integrated reasoning pipelines. `xlog-gpu` provides a high-level deterministic execution API with input/output buffer management; it depends on `xlog-core`, `xlog-cuda`, `xlog-ir`, `xlog-logic`, and `xlog-runtime`. `xlog-prob` implements the full probabilistic tier—provenance tracking, PIR-to-CNF encoding, GPU-native D4 compilation, XGCF circuit construction, exact inference, Monte Carlo sampling, and circuit caching—depending on `xlog-core`, `xlog-cuda`, `xlog-ir`, `xlog-logic`, `xlog-runtime`, and `xlog-solve`. `xlog-induce` is a support subsystem for bounded exact induction on top of the runtime/CUDA stack; it depends on `xlog-core`, `xlog-cuda`, and `xlog-runtime`, and is consumed by the Python ILP surface.

Tier 4 presents user-facing interfaces. `pyxlog` is a PyO3 extension module that bridges nine internal crates (`xlog-core`, `xlog-cuda`, `xlog-gpu`, `xlog-prob`, `xlog-logic`, `xlog-neural`, `xlog-runtime`, `xlog-ir`, `xlog-induce`) to Python, exposing DLPack-first evaluation, training APIs, and ILP/dILP trainers. `xlog-cli` is the published Cargo package that provides the `xlog` command-line binary for deterministic and probabilistic execution with Arrow IPC I/O.

Outside the product/runtime stack, `xlog-cuda-tests` is an unpublished certification suite for GPU release gating, and `xlog-integration` hosts workspace-only cross-crate executor tests and integration binaries.

2.2 Compilation Pipeline

The compilation pipeline transforms Datalog source text into a GPU-executable plan through five stages: parsing, stratification, lowering, optimization, and plan assembly.

Parsing. The PEG parser (built with `pest`) produces an abstract syntax tree of `Program`, `Rule`, `Atom`, `Term`, and `BodyLiteral` nodes. It handles the full xlog language surface (Section 3), including probabilistic facts, annotated disjunctions, neural predicate declarations, user-defined functions, and module imports.

Stratification. The stratifier builds a predicate dependency graph with three edge types—positive, negative, and aggregate—then computes strongly connected components (SCCs) using Tarjan’s algorithm. If any SCC contains a negative or aggregate edge, the program is rejected with a diagnostic error: stratified negation and stratified aggregation require that no recursion passes through negation or aggregation. The output is an ordered sequence of strata, each containing one or more SCCs that can be evaluated before any stratum that depends on them.

Lowering. The lowerer translates each AST rule into a `RirNode` tree. Body literals become `Scan` nodes; variable bindings across atoms produce `Join` nodes with key columns derived from shared variable positions; negated literals become `Anti` joins; filters and arithmetic constraints become `Filter` nodes with `Expr` predicates (supporting comparison, Boolean connectives, arithmetic, built-in functions, and conditional expressions). The lowerer also emits `Fixpoint` nodes that wrap recursive SCC bodies for semi-naive iteration, and `GroupBy` nodes for stratified aggregation.

Optimization. The optimizer applies cost-aware transformations driven by runtime statistics: predicate pushdown and cost-based join ordering, plus promotion of eligible multiway rules to the worst-case-optimal join (WCOJ) subsystem and magic-set rewriting of bound recursive queries. Section 4 details the cost model and the v0.8.6 runtime optimizations (common-subexpression elimination, adaptive re-optimization).

Plan assembly. The final output is an `ExecutionPlan` holding SCCs in dependency order, strata for negation ordering, and compiled rule trees grouped by SCC. The runtime executor interprets it by dispatching each `RirNode` to the appropriate CUDA kernel.

2.3 GPU Residency Model

xlog enforces a hard guarantee: all runtime semantic state resides in GPU device memory, or the system returns a deterministic error. There is no silent out-of-core spilling, and there is no implicit CPU fallback. If a workload exceeds the GPU memory budget, xlog raises a `RESOURCE_EXHAUSTED` error with diagnostic information rather than degrading to host memory transparently.

The residency contract extends to production query paths, where the system targets zero device-to-host (D2H) transfers. The kernel provider accounts for transfers at byte granularity through atomic counters, so a performance-critical section can be bracketed and asserted to perform no column downloads—a check the ILP trainer relies on to prove its credit path stays on-device.

The payoff is bandwidth asymmetry. PCIe 4.0 $\times 16$ delivers roughly 32 GB/s peak bidirectional bandwidth, while GPU HBM provides 1–3 TB/s. A single unnecessary D2H round-trip for a moderate relation (10M tuples at 16 bytes, roughly 160 MB) costs about 5 ms of bus latency—enough to dominate a semi-naive iteration that completes in microseconds on-device. Keeping fact stores, delta relations, circuit values, solver state, and gradient buffers exclusively on the GPU eliminates this class of bottleneck; host involvement is restricted to orchestration, I/O, and compilation.

This accounting is enforced, not advisory: since v0.5.5 a strict D2H guard rejects unexpected data-plane transfers in deterministic evaluation, with a single carve-out for binary-join output counts treated as control-plane metadata. Underneath, a stream-safe device runtime and a **LaunchRecorder** preflight/commit discipline (v0.6.0) keep the recorded paths correct under cross-stream concurrency; later releases build the WCOJ subsystem and the v0.8.6 runtime-optimization layer on the same provider stack (Section 4). Details: `docs/architecture/device-runtime.md` and `recorded-launch-migration.md`.

The **MemoryBudget** configuration in **xlog-core** already carries an `allow_ooc` flag, but current **main** treats it as reserved configuration rather than an implemented spill path. In today’s runtime there is no silent out-of-core execution: workloads that exceed the GPU memory budget fail deterministically with `RESOURCE_EXHAUSTED`.

2.4 IR Stack

xlog uses a layered intermediate representation stack designed for extensibility across reasoning paradigms. Each IR level serves a distinct role in the compilation and evaluation pipeline.

AST (Abstract Syntax Tree). The parser produces a concrete syntax tree that preserves source-level structure: rules, atoms, terms, probabilistic annotations, neural predicate declarations, and directives. The AST lives in `crates/xlog-logic/src/ast.rs` and carries no execution semantics.

RIR (Relational IR). The lowerer transforms the AST into an algebraic tree of **RirNode** variants defined in `crates/xlog-ir/src/rir.rs`. The key variants are: **Scan** (base relation access), **Filter** (predicate evaluation), **Project** (column selection and computation via **ProjectExpr**), **Join** (inner, left-outer, semi, and anti joins with explicit key columns), **GroupBy** (stratified aggregation), **Union**, **Distinct**, **Diff** (set difference for semi-naive delta computation), **Fixpoint** (recursive SCC wrapper), and **MultiWayJoin/ChainJoin** (worst-case-optimal multiway and chain joins promoted from eligible binary-join sequences). RIR nodes carry metadata including estimated cardinality ranges, memory peak estimates, skew hints for join partitioning, and incremental update semantics (delta vs. full materialization).

PIR (Provenance IR). For probabilistic programs, **xlog-prob** constructs a **PirGraph** (defined in `crates/xlog-prob/src/pir.rs`) whose **PirNode** variants—**Lit**, **NegLit**, **And**, **Or**, **Decision**—represent the weighted Boolean formula derived from provenance tracking over the ground program. PIR captures the probabilistic structure needed for knowledge compilation without encoding execution order.

XGCF (xlog GPU Circuit Format). The final compiled form is a leveled DAG stored in `crates/xlog-prob/src/xgcf.rs` as an **Xgcf** structure. Each node has a type (**Const0**, **Const1**, **Lit**, **And**, **Or**, **Decision**), child indices, and variable/literal identifiers. Level offsets enable parallel evaluation: all nodes at the same topological level execute in a single kernel launch. Forward passes compute log-space values; backward passes propagate adjoints for gradient computation. The **GpuXgcf** structure in `crates/xlog-prob/src/gpu.rs` holds the device-resident buffers. Multiple circuits can be batched via `circuit_offsets` for fused evaluation across neural batch dimensions.

Extensibility. The IR stack accommodates additional representations as the system grows. In v0.9.x, **EIR** (Epistemic IR) is the accepted epistemic boundary for world-view reasoning with modal operators, split plans, executable GPU plans, and world-view candidates. **SIR** (Solver IR)

remains the planned Boolean-satisfiability representation for CNF clauses, cardinality constraints, weight vectors, and proof policies. The layered design ensures that adding or extending an IR requires implementing a lowering pass from the appropriate source IR and a backend that targets either XGCF or a GPU-resident evaluation format, without modifying unrelated IR definitions or execution paths.

3 The xlog Language

This section describes the xlog language surface: types, expressions, user-defined functions, modules, and stratified aggregations with constraints. Probabilistic facts (`p::f.`) and neural predicate declarations (`nn/k`) are language-level constructs described in Section 5 and Section 6 respectively. Section 3.7 covers the language-completeness surface added in v0.8.5—finite lists, safe meta-predicates, explicit negation-as-failure, magic-set planning, and probabilistic aggregates.

3.1 Design Philosophy

xlog is a typed logic programming language designed for GPU execution. Three principles shape its surface:

Typed predicates. Every predicate has a declared signature over a closed set of scalar types (`u32`, `u64`, `i32`, `i64`, `f32`, `f64`, `bool`, `symbol`). On current `main`, predicate schemas and lowering-time validation catch argument mismatches before any GPU kernel is generated; richer standalone type inference exists in reserved frontend APIs but is not yet the universal production entry point.

Stratified semantics as contract. Negation, aggregation, and recursion interact through a strongly-connected-component analysis on the predicate dependency graph. Programs whose SCCs cross negation or aggregation edges are rejected at compile time. Every accepted program has a well-defined fixpoint and a deterministic stratum order, which the runtime relies on for GPU scheduling.

One language for four paradigms. Probabilistic facts (`p::f.`), annotated disjunctions, neural predicate declarations (`nn/k`), and SAT constraints share the same syntactic core with deterministic Datalog. Parsing, AST construction, dependency analysis, and most lowering logic are shared across the workspace, while backend-specific expansion and compilation passes refine the resulting plans for deterministic, probabilistic, SAT, and neural-symbolic execution. The distinction between reasoning paradigms lives primarily in the backend, not in the surface syntax.

These three principles are deliberate enablers, not concessions to the GPU. Static types and finite domains are what permit bounded, statically-sized allocation and columnar kernels; stratification is what guarantees a unique fixpoint computable by a deterministic semi-naive schedule. A language admitting unbounded terms or recursion through negation would forfeit exactly the allocation and fixpoint guarantees the GPU runtime depends on—so xlog treats the restriction to finite, stratified programs as the price of compiling logic to efficient device kernels, and Section 3.7 shows how far the finite surface can be pushed before that price is paid.

A minimal xlog program exhibits the core surface:

```
// The minimal xlog program: typed predicates, facts, a recursive rule, a query.
pred edge(u32, u32).
pred reach(u32, u32).
```

```

edge(1, 2). edge(2, 3). edge(3, 4).

reach(X, Y) :-edge(X, Y).
reach(X, Z) :-reach(X, Y), edge(Y, Z).

?-reach(1, N).

```

Listing 1: Minimal xlog program: typed predicates, facts, a recursive rule, a query.

Typed predicate declarations, ground facts, recursive rules, and a query appear in eight lines. The remainder of this section expands each component and adds the extensions that move xlog beyond classical Datalog.

3.2 Type System

Every predicate is declared with a typed signature via `pred`. Eight scalar types are supported: `u32`, `u64`, `i32`, `i64` for fixed-width integers; `f32` and `f64` for floats with IEEE 754 semantics; `bool`; and `symbol` for dictionary-encoded atoms. String literals in source ("`alice`") are terms that map to `symbol` through the `SymbolTable` in `xlog-core`, giving the runtime dense integer identifiers while preserving source-level readability. Users can also introduce named type aliases through `domain` declarations.

Current `main` enforces type consistency through explicit predicate schemas plus frontend/lowering checks. Variable occurrences across body atoms must agree with declared predicate signatures; literal constants are validated against their destination column types; arithmetic expressions preserve scalar types in the emitted `Expr` tree. The dedicated pass in `crates/xlog-logic/src/typeinfer.rs` covers richer inference cases, especially around user-defined functions, but remains a reserved API rather than the sole wired frontend path.

A well-typed program:

```

pred node(u32, symbol).
pred connected(u32, u32).
pred bridge(u32, u32).

node(1, "alice"). node(2, "bob").
connected(1, 2).

bridge(A, B) :-connected(A, B), node(A, _), node(B, _).

```

Listing 2: A well-typed xlog program: typed bridge predicate.

Changing `bridge`'s declaration to `pred bridge(symbol, u32).` would cause inference to report that the head argument `A` is constrained at `symbol` (head) and `u32` (through `connected`), rejecting the program at compile time.

3.3 Expressions and Arithmetic

Rule bodies may contain arithmetic expressions, built-in function calls, and comparison filters. Arithmetic bindings use the `is` operator following Prolog convention: `D is X + Y` computes `X + Y`

and binds the result to `D`. The right-hand side supports the standard arithmetic operators (`+`, `-`, `*`, `/`, `%`), the built-ins `abs`, `min`, `max`, `pow`, and `cast`, and calls to user-defined functions (Section 3.4). Comparison literals (`<`, `<=`, `=`, `==`, `!=`, `>=`, `>`) compile to `Filter` RIR nodes and push down below joins where profitable.

Float arithmetic follows IEEE 754 with explicit NaN and $\pm\infty$ handling: `NaN != NaN`, and ordering predicates place NaN outside any other value rather than raising errors. This matters for aggregations, where a single unguarded NaN would contaminate downstream GPU computations.

A pairwise-proximity query illustrates expression use:

```
pred point(u32, f64, f64).
pred close(u32, u32).

point(1, 0.0, 0.0). point(2, 0.5, 0.5). point(3, 5.0, 5.0).

close(A, B) :-
  point(A, Xa, Ya),
  point(B, Xb, Yb),
  A < B,
  D is (Xa - Xb) * (Xa - Xb) + (Ya - Yb) * (Ya - Yb),
  D < 1.0.
```

Listing 3: Arithmetic in rule bodies: pairwise proximity via squared Euclidean distance.

The rule compiles to a self-join of `point` with an intermediate `Project` that materializes `D`, followed by a `Filter` on `D < 1.0`. The optimizer relocates the cheap comparison `A < B` ahead of the join when profitable, reducing intermediate cardinality.

3.4 User-Defined Functions

xlog includes a user-defined function (UDF) surface for pure scalar computations inside rule bodies and expressions. A UDF is introduced with the `func` keyword; parameters carry optional type annotations, and the return type is declared after a `->`. Arithmetic-expression bodies (optionally with conditional `if/then/else`) are the path integrated into the current arithmetic frontend. The syntax also admits a more ambitious predicate-based form `V :- body.`, but that form remains partially wired and should be viewed as reserved rather than uniformly available across all compiler entry points.

UDFs are intended to be pure: they take scalar arguments and return scalar values, and do not recurse through relations. The function registry in `crates/xlog-logic/src/function.rs` performs structural validation—duplicate definitions, unresolved calls, name conflicts, and recursive definitions without a guarded base case are rejected. Calls to arithmetic UDFs appear wherever arithmetic expressions do, including as the right-hand side of an `is`-expression. Expansion is explicit rather than universal on current `main`: selected execution paths invoke `crates/xlog-logic/src/expand.rs` to inline arithmetic UDFs into composed `Expr` trees (`Add`, `Mul`, `Conditional`, and so on), while predicate-bodied expansion remains reserved. Where this expansion path is used, UDFs participate in the same predicate-pushdown and expression-level optimization passes as built-in arithmetic.

Most classical Datalog dialects have no UDF analogue: functors are either absent or restricted to uninterpreted term construction. xlog’s arithmetic UDF path closes part of this gap without

breaking monotonicity. Because the integrated bodies are pure scalar expressions, a UDF call in a recursive rule does not enlarge the set of derivable facts beyond what the surrounding Datalog rule already derives.

A UDF for score normalization:

```
pred raw_score(u32, f64).
pred norm_score(u32, f64).

func clamp01(X: f64) ->f64 =
  if X < 0.0 then 0.0
  else if X > 1.0 then 1.0
  else X.

raw_score(1, -0.3). raw_score(2, 0.7). raw_score(3, 1.4).

norm_score(Id, N) :-raw_score(Id, R), N is clamp01(R).
```

Listing 4: User-defined function: `clamp01` normalizes a raw score into the unit interval.

The `clamp01` UDF lowers to a conditional expression embedded in the `Project` computing `norm_score`’s second column; no host dispatch occurs during evaluation.

3.5 Modules and Imports

`xlog` programs decompose into modules. A module is a `.xlog` file located on the module search path, which the `xlog` CLI controls via the `-module-path` flag. The `use` directive loads a named module (e.g., `use graph.`) or selectively imports names from it (`use utils/math::{abs, clamp}.`). Module paths use `/` as a separator and mirror the on-disk directory structure. Declarations and functions carry an optional `private` modifier to hide them from importers.

Name resolution runs before type inference. The resolver traverses the import graph, detects cycles, merges symbol tables, and surfaces a single logical program whose remaining compilation stages proceed as if the user had written all rules inline. Imports are transitive; attempts to reach a `private` name yield a structured `ModuleError::PrivatePredicate` diagnostic rather than a silent shadowing.

A two-module program where a top-level entry imports a graph-reachability library:

```
// library: graph_lib.xlog
pred edge(u32, u32).
pred reach(u32, u32).

edge(1, 2). edge(2, 3). edge(3, 4).

reach(X, Y) :-edge(X, Y).
reach(X, Z) :-reach(X, Y), edge(Y, Z).
```

Listing 5: Module `graph_lib.xlog`: predicate declarations and recursive reach rules.

```
// entry: graph_main.xlog
use graph_lib.
```

```
?-reach(1, N).
```

Listing 6: Entry `graph_main.xlog`: imports `graph_lib` and queries `reach`.

The resolver merges the two files before stratification, so the library’s recursive `reach` rules participate in exactly the same SCC analysis and GPU semi-naive evaluation as if they had been written inline.

3.6 Aggregations and Constraints

`xlog` supports stratified aggregation and integrity constraints. Aggregation operators—`count`, `sum`, `min`, `max`, and `logsumexp`—appear at head positions of rules, with an explicit aggregation variable. The stratifier rejects any program in which an aggregation participates in a recursive SCC, guaranteeing a well-defined fixpoint order. Aggregations lower to `GroupBy` RIR nodes, executed on the GPU as radix-sort-and-reduce kernels in `xlog-cuda`.

Integrity constraints are headless rules: `:- body.` denotes a program-level assertion that `body` must be unsatisfiable after evaluation reaches fixpoint. A satisfied constraint is silent; a violated constraint raises a diagnostic. Constraints compose cleanly with stratification—they evaluate after the strata their body predicates inhabit.

A graph-degree computation with an integrity check that every edge source has a degree record:

```
pred edge(u32, u32).
pred degree(u32, u32).

edge(1, 2). edge(1, 3). edge(1, 4). edge(2, 3).

degree(X, count(Y)) :-edge(X, Y).

:-edge(X, _), not degree(X, _).

?-degree(X, N).
```

Listing 7: Stratified aggregation: per-source out-degree with an integrity constraint.

The `degree` rule lowers to a `GroupBy` over `edge` with `count` as the aggregator. The integrity constraint is desugared into an auxiliary rule whose body mirrors the constraint body (so the `not degree(X, _)` literal here contributes an `Anti-join`); at evaluation completion the runtime asserts the auxiliary relation is empty and raises a diagnostic otherwise. `logsumexp` is included explicitly because it is the workhorse aggregation for probabilistic circuits (Section 5), where log-space numerical stability is essential.

3.7 Language Completeness over Finite Domains

Version 0.8.5 expands the surface toward language completeness while preserving the GPU-native execution discipline: every new construct must normalize into the existing typed AST, RIR, probabilistic IR, optimizer, and runtime paths. None of these features introduces a CPU Prolog term heap, dynamic predicate calls, or a hidden CPU-only fallback for a GPU-claimed path; forms that

cannot be lowered safely are rejected at compile time with typed diagnostics rather than executed slowly.

Finite lists. List literals (`[]`, `[A, B]`, cons patterns `[H|T]`) and the `list<T>` column type denote finite, homogeneous, typed lists. Lists and the core list built-ins normalize to `__xlog_list_*` helper relations that lower through ordinary relational joins and aggregations; open-ended generators, cyclic lists, and untyped heterogeneous lists are rejected.

Safe meta-predicates. Finite, statically-resolvable forms of `ground`, `var`, `nonvar`, `functor`, `=.. (univ)`, source-fact `findall`, and unary/binary `maplist` normalize to `__xlog_meta_*` helper relations before stratification. Term, compound, and `predref` columns lower to u64 term identifiers. Unrestricted `call/N`, dynamic database mutation, and runtime-variable predicate names remain outside the contract.

Explicit negation-as-failure. Deterministic `not atom(...)` is closed-world stratified negation over deterministic relations, lowering to the existing anti-join path. Source-order binding-safety diagnostics require every named variable in a negated atom to be bound by a prior positive binder. Probabilistic well-founded semantics (Section 5) remains a separate exact-inference profile.

Magic-set planning. The `#pragma magic_sets = auto|on|off` directive enables adornment-based rewriting of bound deterministic recursive queries: a query such as `?- reach(1, Y).` specializes recursive reachability to the reachable frontier, reducing materialized rows. The rewrite runs before optimizer and WCOJ planning and declines—with a typed diagnostic—when equivalence cannot be proven (e.g., unsafe interaction with negation, aggregates, or mutual recursion).

Probabilistic aggregates. Finite `count`, `sum`, `min`, `max`, and `logsumexp` outcomes are supported as aggregate heads in exact provenance/PIR inference and in Monte Carlo deterministic aggregate execution, with typed exact-domain cap diagnostics. Small finite `count` domains use aggregate lifting—exact cardinality dynamic programming over a compact state set instead of naive world enumeration—and accepted fired `count` aggregates dispatch the GPU-native `weights_count_lift_exact` evaluator rather than a CPU finite-world shortcut.

Approximate inference and incremental parsing. Monte Carlo configuration is a first-class language feature via `#pragma prob_engine = mc` plus `sample/seed/confidence` pragmas, with documented source/CLI precedence and reported uncertainty. A statement-level incremental parser with stable spans reuses unchanged parses across edits, supporting interactive `xlog repl` and `xlog watch` developer workflows.

4 GPU-Native Datalog Evaluation

Section 3 describes the source language; this section describes how its Datalog subset compiles to and executes on the GPU. The deterministic evaluation engine executes standard Datalog with stratified negation and aggregation entirely on the GPU. The algorithmic approach—semi-naive fixpoint iteration over stratified programs—is well established. The contribution here is not algorithmic novelty but engineering: every relational operator runs as a custom CUDA kernel, delta relations are maintained on-device, and no host–device transfers occur during evaluation.

4.1 Semi-Naive Evaluation on GPU

The executor in `crates/xlog-runtime/src/executor/` processes an `ExecutionPlan` consisting of strata ordered by the dependency analysis from Section 2. Each stratum contains one or more strongly connected components (SCCs), and the executor dispatches each SCC according to whether it is recursive or non-recursive.

Non-recursive SCCs. The executor evaluates each compiled rule once, passing the RIR tree to the node dispatcher which invokes the appropriate CUDA kernel for each operator (scan, join, filter, project, group-by). Results for the same head predicate are merged via the `union_gpu` kernel and deduplicated. No iteration is needed.

Recursive SCCs. The `execute_recursive_scc` method in `crates/xlog-runtime/src/executor/recursive.rs` implements semi-naive evaluation. The algorithm proceeds in three phases:

1. **Seeding.** Every rule in the SCC is evaluated once against the current relation store, and each predicate’s initial delta is the set difference between the newly derived tuples and any pre-existing ones. Delta relations are distinct on-device buffers.
2. **Iteration.** Each rule is re-evaluated through delta-rewritten variants—one variant per recursive scan occurrence, so self-joins are handled correctly: a body containing $p(X,Y)$, $p(Y,Z)$ yields two variants, each substituting the delta into exactly one scan. Variant results are unioned per head predicate, differenced against the full relation, and deduplicated, all on-device. When no predicate produces new tuples, the fixpoint is reached.
3. **Merge and cleanup.** New delta tuples are merged into the full relation by union and sorted deduplication. At fixpoint—or when the iteration cap (default 1,000,000) is hit—delta relations are freed.

All intermediate buffers remain on-device throughout. The profiler records per-operator timing, row counts, and peak memory, feeding the optimizer loop described below.

4.2 Kernel Design

xlog implements its relational algebra as eight direct CUDA kernels rather than wrappers over cuBLAS, cuSPARSE, or Thrust. The reason is the residency contract: a library operator imposes its own allocation, synchronization, and host-staging behavior, which would defeat the zero-transfer guarantee of Section 2. Bespoke kernels keep every intermediate buffer on-device and let every launch flow through the recorded-launch discipline. Table 1 summarizes them; the worst-case-optimal multiway join described below is a ninth kernel family.

Every operator also exposes a recorded-launch sibling (since v0.6.0) that routes through the `LaunchRecorder` preflight/commit discipline of Section 2, selectable per-operator or via an umbrella flag, with the legacy path as the default.

IEEE 754 total ordering. The filter kernel implements a correctness-critical design choice for floating-point comparisons. Equality and inequality (Eq, Ne) use standard IEEE 754 semantics where $\text{NaN} \neq \text{NaN}$. Ordered comparisons (Lt, Le, Gt, Ge) use a total ordering transformation: the `float_to_ordered_f64` device function converts an `f64` bit pattern to an `i64` such that the resulting integer

Table 1: Core relational CUDA kernels in `xlog-cuda`.

Kernel	Purpose
<code>join.cu</code>	Hash join with linked-list collision chains and FNV-1a composite hashing over raw key bytes, so all scalar types share one path. Includes count $\rightarrow$ scan $\rightarrow$ materialize (CSM) variants for Inner and LeftOuter joins.
<code>sort.cu</code>	Stable 4-bit radix sort (8 passes over 32-bit keys): per-pass histogram, prefix sum, scatter.
<code>filter.cu</code>	Comparison operators for column-vs-constant and column-vs-column, with stream compaction via prefix sum and fused compare-and-scan variants.
<code>groupby.cu</code>	Sorted-input grouped aggregation: boundary detection in sorted keys, then per-group Count/Sum/Min/Max/LogSumExp reduction.
<code>dedup.cu</code>	Sort-based deduplication and set difference with type-aware equality, including IEEE 754 float handling.
<code>set_ops.cu</code>	Union (concat + sort + dedup) and set difference via sorted-array binary search.
<code>scan.cu</code>	Blelloch parallel prefix sum—the building block for filter compaction, dedup, and group-by.
<code>pack.cu</code>	Multi-column key packing into row-major byte arrays with FNV-1a hashing, eliminating host roundtrips for join-key preparation.

order matches the IEEE 754 `totalOrder` predicate. The ordering is: $-\text{NaN} < -\text{Inf} < \text{negative finites} < -0.0 < +0.0 < \text{positive finites} < +\text{Inf} < +\text{NaN}$. The transformation works by flipping all bits except the sign bit for negative values and flipping only the sign bit for non-negative values, matching Rust’s `f64::total_cmp()` algorithm. An analogous `float_to_ordered_f32` function handles 32-bit floats. This ensures that Datalog programs produce deterministic, well-defined results for all floating-point inputs, including edge cases involving NaN and signed zeros.

4.3 Adaptive Join Planning

The optimizer performs cost-based query transformations using runtime statistics collected during prior executions.

Cost model. Each plan node is costed along four dimensions—expected row count, CPU coordination cost, peak GPU memory, and number of host–device transfers—with transfers penalized heavily to reflect PCIe latency. The transfer penalty is what makes the planner prefer device-resident plans even when they cost more raw compute.

Predicate pushdown. When enabled (the default), the optimizer pushes **Filter** nodes below **Project** and **Join** nodes. Filters above joins are decomposed: predicates referencing only left-side columns are pushed into the left input, predicates referencing only right-side columns into the right input, and cross-predicates remain above the join. Filters above projections are remapped through pass-through columns where possible.

Join ordering. For queries of up to ten relations the optimizer uses dynamic programming to select the minimum-cost join order; above that threshold it falls back to a greedy heuristic to avoid exponential planning time. Cardinality estimates multiply left and right cardinalities by a selectivity factor learned from prior executions.

Feedback loop. The executor can export a statistics snapshot—per-relation cardinality, byte size, access heat, column statistics, and join selectivities—and merge it back before the next compilation, closing the loop between execution and optimization. A predicate-name mapping keeps these statistics valid across compilations even when internal relation IDs are reassigned.

4.4 Worst-Case Optimal Joins

Binary-join plans can materialize intermediate relations that are asymptotically larger than the final result—the classic pathology for cyclic queries such as triangle enumeration. This is not a corner case for the Datalog workloads xlog targets: program-analysis queries (points-to analysis, call-graph and reachability inference over code-dependency graphs) are dominated by cyclic, skewed join patterns where the intermediate blowup, not the final answer, governs cost. xlog addresses this with a worst-case-optimal join (WCOJ) subsystem [12–14]. A pure-Rust hypergraph planner (`xlog-logic::hypergraph`) analyzes rule bodies for multiway-join eligibility, derives a deterministic variable ordering from cardinality and selectivity statistics, and emits a `MultiWayJoin/ChainJoin` RIR node when promotion provably preserves output semantics; otherwise the binary-join plan is retained as a certified fallback. Planner verdicts and GPU outputs are certified against the same CPU oracle fixtures before executor wiring.

The GPU side (`wcoj.cu`, 2,262 LOC) executes the join via `count` $\rightarrow$ device-side prefix scan $\rightarrow$ materialize over flat sorted-column storage, with a four-byte metadata D2H total and no count-vector transfer. A sorted-layout fast path proves already-sorted-and-unique inputs with a recorded checker kernel and clones on-device instead of re-sorting. Coverage extends from the canonical triangle to 4-cycles, K-clique templates (K=5–8), recursive/SCC integration, and helper-relation splitting that exposes inner skew to root-level histograms.

A default-on adaptive skew classifier decides per query whether to dispatch WCOJ: high-skew inputs take the WCOJ path while uniform or empty inputs fall back to the binary-join chain. `XLOG_DISABLE_WCOJ_TRIANGLE=1` is the hard kill switch. Section 9 reports the measured speedup envelope on skewed paper-scale workloads.

4.5 Runtime Optimization Layer

Interactive and long-running deployments—REPL sessions, iterative solver loops, incremental ML training—repeatedly evaluate similar queries over relations that change little between calls. A stateless executor rebuilds join indices and recomputes shared subplans on every call, paying the dominant cost each time. v0.8.6 adds three runtime optimizations that exploit this between-call stability, each with explicit off/on controls and zero added data-plane host transfers. *Common-subexpression elimination* caches the results of safe deterministic subplans within an evaluation, invalidated on relation-generation change, removing duplicate work. *Adaptive re-optimization* lets the executor adopt a compiler-supplied candidate plan (via `execute_plan_with_adaptive_candidate`) when deterministic mis-plan telemetry crosses a configurable ratio, with GPU output equivalence and rollback diagnostics. A *persistent hash-index manager*, keyed on relation generation, schema, and device, retains built indices across repeated session evaluations with deterministic LRU eviction and background-build accounting; on a build-heavy repeated-session fixture it records a $3.21\times$ speedup against a $1.5\times$ target with zero tracked host transfers (Section 9). Persistent sessions additionally coalesce relation-delta batches and expose opt-in, metadata-only relation-change callbacks.

4.6 Reversible Symbols

Datalog programs operate on strings, but GPU kernels operate on fixed-width numeric columns. xlog bridges this with a global symbol table that interns each unique string to a dense, sequential

`u32` identifier and resolves it back in $O(1)$ at output time. Dense zero-based IDs give GPU-friendly access patterns, and symbol columns then *are* ordinary `u32` columns—join, sort, filter, and dedup kernels operate on them with no special-casing. The table itself stays host-side, since variable-length string data is needed only at ingestion and output boundaries. The `symbol` type maps to Arrow’s `Dictionary(UInt32, Utf8)`, so export to Arrow consumers (cuDF, Polars, DuckDB) is zero-copy while preserving the compact internal representation.

5 Probabilistic Inference on GPU

5.1 Approach

Probabilistic logic programming systems such as ProbLog [15] compute the probability of queries over programs containing probabilistic facts. The standard technique is knowledge compilation: compile the ground Boolean formula derived from provenance tracking into a tractable circuit (typically a Decision-DNNF), then evaluate weighted model counts over that circuit. ProbLog follows a compile-once/evaluate-many model: the circuit structure is determined by the logical program and is independent of the probability weights, so a single compilation can serve many weight configurations during parameter learning.

`xlog` adopts the same compile-once/evaluate-many model but moves the entire pipeline onto the GPU. In ProbLog, provenance extraction, CNF encoding, D4 compilation, and circuit evaluation all execute on the CPU. When ProbLog is paired with a neural component (as in DeepProbLog [1]), each training iteration requires transferring circuit evaluation results from CPU to GPU for gradient computation, and pulling gradients back for weight updates. `xlog` eliminates this architectural split: PIR extraction, Tseitin [16] CNF encoding, D4 [7] compilation, CDCL equivalence verification, and forward/backward circuit evaluation all execute as CUDA kernels on the GPU. The compiled circuit is a GPU-resident data structure (XGCF), weight buffers and gradient buffers are GPU-resident, and no host–device transfers occur in the evaluation hot path. This design turns what would be a PCIe-bottlenecked ping-pong between CPU logic engine and GPU tensor framework into a single-device pipeline.

5.2 GPU Knowledge Compilation Pipeline

The probabilistic compilation pipeline transforms a Datalog program with probabilistic annotations into a GPU-resident arithmetic circuit through five stages (Figure 2). Each stage is implemented as one or more CUDA kernel launches, with intermediate data structures remaining in device memory throughout.



Figure 2: Knowledge compilation pipeline. All five stages execute as CUDA kernels on GPU-resident data; no host–device transfers occur between stages.

The five stages map to the kernel files `pir.cu`, `cnf.cu`, `d4.cu`, `sat.cu`, and `circuit.cu`.

Stage 1: PIR extraction. Provenance tracking over deterministic evaluation produces a weighted Boolean formula. The interner hash-conses node batches on the GPU, deduplicating structurally identical subformulas through an atomic GPU hash table. This on-device deduplication is essential: grounding can produce exponentially many redundant subformulas, and interning on the host would force the full uninterned graph into host memory—reintroducing the very transfer the pipeline exists to avoid.

Stage 2: CNF encoding. A Tseitin transformation lowers the formula to CNF entirely on GPU, using parallel prefix sums to assign compact, gap-free variable IDs and emit clauses into a device-resident compressed-sparse-row structure.

Stage 3: D4 compilation. The most complex stage compiles the CNF into a Decision-DNNF circuit (the XGCF format) via a GPU-native variant of the D4 algorithm [7]. A BFS frontier exposes parallelism, then per-block DFS workers perform component detection, decision-variable selection via VSIDS-like activity scores [17], and recursive decomposition. A smoothing pass forces every branch of an OR/Decision node to mention the same variable set—required for correct weighted model counting—and the result is leveled so each topological level evaluates in one kernel launch.

Stage 4: CDCL verification. Rather than trust the compiler, xlog proves the compiled circuit C is equivalent to the source formula φ : it solves $\varphi \wedge \neg C$ and $C \wedge \neg \varphi$ on a GPU CDCL solver, both of which must be unsatisfiable. This is a complete formal proof, not a sampling check. UNSAT results carry a resolution proof checked on-device, and the solver never reads its status back to the host—a wrong answer triggers a device-side trap. Circuit correctness thus becomes a machine-checked certificate rather than an assumption, which matters because a silently miscompiled circuit would corrupt every downstream probability and gradient.

Stage 5: circuit evaluation. The verified XGCF supports a level-parallel forward pass (log-space weighted model counting, one kernel per level, yielding $\log Z$) and a reverse-order backward pass that accumulates per-variable gradients. Both leave their results in GPU-resident buffers consumable directly by PyTorch via DLPack (Section 6).

5.3 Circuit Caching

D4 compilation is the most expensive stage in the pipeline. On the MNIST addition benchmark (01_minimal, 512 training examples), a cold-start compilation takes approximately 75 seconds. Since this cost is incurred at the beginning of each training epoch in a naive implementation, it can dominate total training time.

The key insight enabling circuit caching is that the XGCF circuit structure depends on the logic program—specifically, on the propositional structure of the grounded rules and the evidence pattern—but not on the numerical values of the probability weights. When neural network parameters change between training iterations, the weights assigned to probabilistic facts change, but the Boolean formula and its compiled Decision-DNNF structure remain identical. This means the compiled XGCF can be cached and reused across all training iterations, with only the weight buffers updated before each forward pass.

The `GpuCircuitCache` in `crates/xlog-prob/src/compilation/gpu_cache.rs` implements this caching. On the first evaluation, the full pipeline runs (PIR extraction, CNF encoding, D4 compilation, CDCL verification), and the resulting `GpuXgcf` is stored in the cache keyed by a structural hash of the PIR graph and evidence configuration. On subsequent evaluations, the cache

returns the existing circuit, and only the weight-upload and forward/backward passes execute. The cache also stores the CNF variable tables (`GpuCnfVarTables`) needed to map between PIR leaf IDs and CNF variable IDs for weight construction.

Measured impact: on the MNIST addition benchmark (3 epochs, 512 training examples, 3 seeds), cache-enabled training completes in a mean of 88.89 seconds (std: 3.48s), compared to 242.90 seconds (std: 8.19s) with caching disabled. This yields a $2.74\times$ end-to-end speedup (95% CI: [2.29, 3.18]). The speedup is large because D4 compilation cost is amortized across all epochs: in the latest checked-in post-fix rerun on `main`, the first epoch averages 71.654s while warm epochs average 0.248s once the cached circuit and batched evaluation path are active (Section 9).

5.4 Monte Carlo Sampling

For programs where exact inference via knowledge compilation is infeasible—either because the ground formula is too large for D4 compilation within the available GPU memory, or because the user prefers approximate results with bounded computation—xlog provides a GPU-parallel Monte Carlo sampling engine implemented in `crates/xlog-prob/src/mc/`.

The MC engine supports two sampling methods, selectable via `McSamplingMethod`:

- **Rejection sampling.** The sampler draws values for all probabilistic facts from their prior distributions on the GPU, then evaluates the deterministic Datalog program in each sampled world. Worlds where the evidence is not satisfied are discarded. Query probabilities are estimated as the fraction of accepted worlds in which each query atom is derived. This method is unbiased but can be inefficient when the evidence has low prior probability, as most samples are rejected.
- **Evidence clamping.** When all evidence atoms correspond to Bernoulli probabilistic facts, the sampler can force those facts to their observed values, guaranteeing that every sample satisfies the evidence. This eliminates rejection waste entirely: the `McCountStrategy` switches to `QueriesOnly` mode, skipping evidence-side buffer allocations and truth-kernel evidence checks. Evidence clamping is auto-selected when the `EvidenceForcing` analysis determines that all evidence is forceable (i.e., each evidence atom maps to a single probabilistic fact with no intermediate rules).

Both methods report confidence intervals for estimated probabilities. The sampling loop is structured to minimize per-sample overhead: probabilistic fact values are sampled in bulk on the GPU via the `sampling` submodule, the relation store is reset using a targeted `McSampleResetPlan` that preserves pure deterministic relations and clears only dynamic relations (avoiding a full store clone per sample), and per-sample pointer buffers are pre-allocated outside the sample loop. The current repository contains these hot-loop optimizations in code, but does not ship a standalone archived MC timing report with stable phase-by-phase numbers, so this paper limits itself to the implementation strategy rather than quoting unsupported per-benchmark deltas.

5.5 Probabilistic Aggregates

v0.8.5 extends probabilistic inference with finite aggregate queries. Aggregate heads and outputs over `count`, `sum`, `min`, `max`, and `logsumexp` are evaluated in both exact provenance/PIR inference

and Monte Carlo deterministic aggregate execution, subject to a typed exact-domain cap that rejects domains too large for tractable enumeration (with a diagnostic recommending MC or a smaller finite domain). For small finite `count` domains, *aggregate lifting* replaces naive world enumeration with exact cardinality dynamic programming over a compact state set—collapsing, for example, a 2^{17} -outcome enumeration into roughly 170 lifted states with identical results. Accepted fired `count` aggregate queries dispatch a GPU-native exact evaluator (`weights_count_lift_exact`) rather than a CPU finite-world shortcut, keeping the exact aggregate path device-resident; explain metadata records which aggregate operators fired versus fell back to exact enumeration.

5.6 Well-Founded Semantics

Standard stratified negation requires that no recursive cycle passes through a negated literal. Programs that violate this constraint—such as the classic `p :- not q. q :- not p.`—have no two-valued stable model and are rejected by most Datalog engines. `xlog` handles such non-monotone programs via Well-Founded Semantics (WFS), implemented in `crates/xlog-prob/src/wfs.rs`.

WFS assigns three-valued truth to atoms: true, false, or undefined. The algorithm computes an alternating fixed point: (1) find the greatest unfounded set (atoms that cannot be supported by any rule), mark them false; (2) derive all consequences of the current knowledge, mark them true; (3) repeat until stable. Atoms that remain neither true nor false are classified as undefined. For probabilistic programs, true atoms receive normal probability and gradient computation, while false and undefined atoms are assigned probability zero with zero gradients—matching ProbLog’s conservative treatment of non-stratified programs.

WFS is invoked automatically during provenance extraction when a non-monotone SCC is detected. The Monte Carlo engine, by contrast, is the GPU-resident Datalog megakernel (`crates/xlog-prob/src/mc/resident.rs`): it evaluates every sampled world in a single device launch with a device-side bounded fixpoint and supports only *bounded-domain positive* Datalog (arity ≤ 2 , body ≤ 2). Non-monotone and otherwise unsupported fragments are *rejected fail-closed* before execution with a typed diagnostic (there is no host-orchestrated fallback); such programs must use the exact / WFS path. A representative supported MC program is `examples/prob/04-recursive-mc.xlog`, a recursive transitive closure (`reach(X,Z) :- reach(X,Y), edge(Y,Z).`) whose non-base tuple `reach(1,3)` is derived by a multi-iteration device-side fixpoint.

6 Neural-Symbolic Bridge

The previous sections describe `xlog` as a purely symbolic system: Datalog evaluation on the GPU (Section 4), probabilistic inference through knowledge compilation (Section 5). This section introduces the neural-symbolic bridge—the machinery that connects PyTorch neural networks to the probabilistic logic engine, enabling end-to-end differentiable learning where neural perception and logical reasoning are jointly trained.

6.1 Neural Predicates

In `xlog`, a neural network *is* a predicate. The `nn/4` declaration embeds a network directly into the Datalog program:


```
nn(network_name, [input_vars], output_var, [output_labels]) ::predicate(args).
```

Listing 8: Neural predicate declaration syntax (`nn/4`).

This is not a foreign-function call or an escape hatch. The declaration states that evaluating predicate `predicate(args)` requires a forward pass through the named network. The network’s softmax output over the label set becomes a distribution over probabilistic facts, which then participates in the standard knowledge compilation pipeline described in Section 5.

On the Python side, networks are registered via the `pyxlog` API. The `CompiledProgram.register_network()` method accepts a PyTorch `nn.Module`, an optimizer, and optional configuration parameters:

```
program.register_network("mnist_net", net, optimizer)
```

Listing 9: Registering a PyTorch module against a neural predicate via `pyxlog`.

Under the hood, the Rust-side network registry manages a handle per network, each holding a reference to the PyTorch module, optimizer, and optional learning-rate scheduler. The `NetworkConfig` struct exposes builder-pattern configuration for batching (grouping multiple queries into a single forward pass), top- k sampling (restricting the output space when most classes have near-zero probability), deterministic mode (argmax instead of probabilistic sampling), and output caching (avoiding redundant forward passes for repeated inputs, up to a configurable cache size of 10,000 entries by default).

The data flow through a neural predicate proceeds in two directions. **Forward:** when the logic engine encounters an `nn/4`-backed predicate during evaluation, it invokes the registered PyTorch module. The network’s softmax output vector is decomposed into weighted probabilistic facts—one per output label—which are fed into the PIR/CNF/D4/XGCF compilation pipeline. Because the circuit structure depends on the logic program and not on the weight values, the compiled XGCF template is cached (Section 5) and only the leaf weights are updated on subsequent iterations. **Backward:** after the XGCF circuit computes the forward log-probability and backward gradients via level-parallel CUDA kernels, the per-leaf gradient buffers are exported as `DLPack` tensors and injected into PyTorch’s autograd graph. The optimizer then updates the network parameters as usual. The Python GIL is released during GPU-resident work using `py.allow_threads()`, ensuring that CUDA kernel execution and GPU memory operations do not block the Python interpreter.

Production neural-bridge surface. For downstream consumer integration, the v0.8.0 bridge exposes registered-network top- k and deterministic output modes with stable top- k tie-breaking, neural-cache hit/miss telemetry, and a family of loss helpers—Belnap pro/contra/quarantine losses for paraconsistent four-valued supervision, semantic loss, mean-squared error, and an information-theoretic loss—so that training pipelines drive the symbolic engine without reaching into internal crates. Native exact induction is reachable from the same Python surface (`pyxlog.ilp.exact_induce.induce_exact(..., backend="native")`), routing tensorized-ILP candidate scoring through the GPU `ilp_exact` kernel with a constant device-to-host transfer budget.

6.2 End-to-End Training Loop

To make the integration concrete, we walk through the MNIST addition example—the canonical DeepProbLog benchmark. The task: given two MNIST digit images, predict their sum. The network never sees individual digit labels; it learns to classify digits purely from supervision on sums.

The logic program. The complete xlog source (`examples/neural/01_minimal/minimal.xlog`) is:

```
// Neural predicate declaration
// nn(network_name, [input_vars], output_var, [output_labels]) ::predicate(args).
//
// This declares that:
// - 'mnist_net' is the name of the neural network
// - [X] means the network takes one input (image at index X)
// - Y is the output variable (the classified digit)
// - [0,1,2,3,4,5,6,7,8,9] are the possible output labels (digit classes)
// - digit(X, Y) is the resulting predicate
//
// The network's softmax output becomes probabilistic facts:
// P(digit(0, 0)) = network_output[0]
// P(digit(0, 1)) = network_output[1]
// ...
// P(digit(0, 9)) = network_output[9]
nn(mnist_net, [X], Y, [0,1,2,3,4,5,6,7,8,9]) ::digit(X, Y).

// Addition rule
// The sum of two digits at indices X and Y is Z.
//
// This uses probabilistic inference:
// P(addition(0, 1, 7)) = sum over all (d1, d2) where d1 + d2 = 7 of:
// P(digit(0, d1)) * P(digit(1, d2))
//
// For example, with d1=3, d2=4:
// P(addition(0, 1, 7)) includes P(digit(0, 3)) * P(digit(1, 4))
addition(X, Y, Z) :-digit(X, D1), digit(Y, D2), Z is D1 + D2.
```

Listing 10: MNIST addition program: a CNN-backed digit predicate plus an addition rule.

Two lines of logic define the entire reasoning structure. The `nn/4` declaration connects the `mnist_net` CNN to a `digit/2` predicate. The addition rule computes all consistent decompositions of a sum through probabilistic marginalization.

The training script. The Python driver (`examples/neural/01_minimal/train.py`) compiles the program, registers the network, and runs training through the `pyxlog.train_model_tensor` API:

```
program = pyxlog.Program.compile("""
    nn(mnist_net, [X], Y, [0,1,2,3,4,5,6,7,8,9]) ::digit(X, Y).
    addition(X, Y, Z) :-digit(X, D1), digit(Y, D2), Z is D1 + D2.
""")

program.register_network("mnist_net", net, optimizer)
```




Figure 3: One iteration of neural-symbolic training on MNIST addition. The XGCF circuit is compiled once (first iteration) and reused thereafter; only the leaf weights change across iterations. Gradients flow back through DLPack capsules into PyTorch’s autograd graph.

```

program.add_tensor_source("train", train_images)

queries = generate_queries(n_pairs, train_labels)

history = pyxlog.train_model_tensor(
    program,
    queries,
    epochs=args.epochs,
    batch_size=args.batch_size,
    log_iter=args.log_iter,
)

```

Listing 11: Python driver for MNIST addition: compile, register network, train.

Each query has the form `addition(i, j, expected_sum)`. The training objective minimizes the negative log-probability of the query under the compiled circuit: $-\log P(\text{addition}(i, j, s))$. The `train_model_tensor` function orchestrates the full loop: batching queries that share the same circuit template, running forward and backward passes over the cached XGCF circuit, and propagating gradients back to the PyTorch optimizer.

Training loop architecture. Figure 3 shows one training iteration.

The critical performance property is that the circuit compilation (D4) happens once, during the first forward pass. All subsequent iterations reuse the cached XGCF structure, executing only the weight update and the level-parallel forward/backward kernels. This is the mechanism behind the $2.74\times$ training speedup reported in Section 5.

Results. On the MNIST addition benchmark (512 training images, 256 addition pairs, batch size 64, learning rate 10^{-3} , CUDA device), the latest checked-in post-fix rerun on `main` reports the summary in Table 2 over 3 seeds:

Table 2: MNIST addition training summary (01_minimal, post-fix rerun, 3 seeds).

Metric	Value	Artifact	Notes
Mean first epoch	71.654 s	Track B comparison	Cold start with D4 compile/verify
Mean steady epoch	0.248 s	Track B comparison	Warm batched path after cache reuse
Mean total train	72.645 s	Track B comparison	5 epochs, 3 seeds
Mean final loss	2.62953	Track B comparison	Final training loss over the 3 recorded seeds

These measurements come from the post-fix Track B comparison artifact at `examples/neural/results/track_b/20260218T145032Z_track_b/comparisons/mnist_xlog_vs_scallop.json`, which exercises the full xlog pipeline—compilation, circuit caching, GPU-resident forward/backward, and batched query evaluation—through `pyxlog.train_model_tensor`. The loss reduction across the recorded runs confirms that gradients flow from circuit evaluation through the logic program back to the CNN parameters, and the timing split matches the cold-start vs. warm-epoch behavior reported in Section 9.

6.3 Term Embeddings (v0.5.0)

Version 0.5.0 introduces a term embedding API that attaches dense vector representations to Datalog symbols. Where `nn/4` neural predicates map perception (images, text) to discrete logical facts, term embeddings go in the opposite direction: they give logical entities a continuous representation that can participate in neural computation.

The Python API provides two methods on `CompiledProgram`:

- **`register_embedding(name, module_or_tensor, trainable=True)`** accepts either a PyTorch `nn.Embedding` module or a raw 2D `torch.Tensor`. For `nn.Embedding`, the weight matrix shape `[vocab_size, dim]` is read from the module. For raw tensors, `trainable` must be `False` (frozen embeddings). The Rust side stores the module reference, dimensionality, vocabulary size, and trainability flag, registered in the same registry as neural predicates and keyed by name.
- **`forward_embedding(name, ids)`** looks up embedding vectors by integer IDs, returning a batched `torch.Tensor` with shape `[len(ids), dim]`. For trainable `nn.Embedding` modules, the returned tensor carries autograd history, so gradients flow through the embedding lookup. For frozen tensors, the result has `requires_grad=False`. Internally, the method determines the device from the embedding weight tensor and creates the ID tensor on the same device to avoid CPU/CUDA mismatches.

This API enables entity representation learning within the logic framework. For example, a knowledge graph completion task can attach learned embeddings to entity symbols, combine them with relational predicates in xlog rules, and jointly train the embeddings alongside neural perception networks. The embeddings are the only gradient-carrying embedding API as of v0.9.2; future releases may integrate embedding lookups directly into the circuit evaluation path.

6.4 Differentiable ILP (Beta)

Note: The differentiable ILP subsystem is in **beta** as of v0.9.2. The API surface, promotion pipeline, and artifact format may change in future releases. The schema version is **beta-v2**.

xlog includes a differentiable inductive logic programming (dILP) trainer [6] that learns first-order rules from examples. Unlike the neural predicate path (which trains network parameters given fixed rules), dILP searches the space of candidate rules, using gradient-based optimization to assign soft probabilities to candidates and hard selection to commit the best rule.

Sparse GPU mask API. The scalability bottleneck in dILP is the candidate mask: for R relations of arity A the naive candidate space is $O(R^A)$, so materializing a dense mask on the Python side is prohibitive. xlog instead offers a graduated family of sparse-mask entry points that trade Python convenience for strictness—from accepting a full DLPack probability vector and selecting top- k on the Rust side, to a fully device-resident path where preselected candidate IDs never leave the GPU (resolved against a fixed, pre-uploaded candidate order). The strictest path keeps the entire mask-construction step transfer-free.

The GPU loss/gradient computation (`compute_ilp_loss_grad_gpu`) builds a sparse COO structure from the per-entry membership masks, converts it to CSR via device-side sort, and launches forward (credit gather + NLL loss) and backward (gradient scatter) CUDA kernels. The entire credit/loss path is GPU-resident with zero host transfers when using the strict sparse device mask.

Six-gate promotion pipeline. When training converges and a candidate rule is selected, xlog does not blindly commit it to the program. Instead, the rule passes through a six-gate promotion pipeline:

1. **Convergence gate**—verifies that training actually converged (loss plateau, discreteness threshold).
2. **Novel-rate audit**—checks the fraction of new derivations the rule produces; a configurable `max_novel_rate` threshold catches rules that overgenerate.
3. **Regression check**—ensures the new rule does not invalidate existing derivations for protected relations.
4. **Holdout F1 gate**—evaluates the candidate on held-out examples using k-fold cross-validation (default 5 folds, threshold 0.95 F1).
5. **Ambiguity scan**—detects when multiple candidates have similar soft probabilities, optionally reporting `ambiguous_alternatives` for manual review.
6. **Typed schema validation**—when `typed_schema_required=True`, verifies that the candidate rule’s variable bindings respect declared type annotations on relations.

A rule that fails any gate receives a `PromotionStatus` of `GATE_FAILED` or `MANUAL_REVIEW_REQUIRED`, with per-gate `GateResult` details. Only rules that pass all six gates are committed.

Hard-negative mining. The `TrainConfig.max_mined_negatives` parameter enables hard-negative mining: the trainer identifies examples that the current candidate rule incorrectly covers

and adds them as negative examples for subsequent training rounds. This improves precision for rules operating in domains with many near-miss patterns.

Artifact persistence. Learned rules are serialized as `LearnedArtifact` objects containing the candidate map, final logits, soft probabilities, hard selection, discovered rule text, full `TrainConfig` snapshot, and training telemetry. The `save()` method writes JSON with SHA-256 hashes computed over both the candidate map and the configuration snapshot, enabling integrity verification on `load(path, verify_hash=True)`. The schema version (`beta-v2`) is validated on load to prevent silent format incompatibilities. For strict GPU-native training paths where logits and probabilities remain device-resident, a `StrictLearnedArtifact` variant defers host materialization until `export_compat_artifact()` is explicitly called, avoiding unnecessary device-to-host transfers during automated pipeline runs.

7 Epistemic Reasoning

The reasoning modes described so far ask what is *true*: which atoms a program derives, with what probability, or under what neural evidence. Epistemic logic programming instead asks what is *known* or *possible* across the multiple stable models a program admits—the natural setting for introspection and reasoning under incomplete knowledge, where a query ranges over a program’s whole space of world views rather than a single model. xlog adds this mode in its v0.9.x line, executing it on the same GPU-resident substrate as the other backends.

7.1 Modal Surface and the Epistemic IR

A program may use four modal literals in rule bodies—`know`, `possible`, `not know`, and `not possible`—over a predicate. Rather than rewrite these away at parse time, the compiler preserves them in an *Epistemic Intermediate Representation* (EIR) that sits between the AST and GPU lowering, so the world-view semantics are resolved by a dedicated execution plan rather than smuggled into ordinary relational algebra. The accepted surface spans finite nested modal-operator chains (normalized and executed), epistemic integrity constraints (which prune candidate world views on the GPU), rules that jointly couple several epistemic predicates across components, same-name multi-arity modal predicates, and recursive epistemic execution (Section 7.3). The boundaries that remain fail-closed are genuine rather than incidental: raw RIR lowering of modal literals, genuinely cyclic modal coupling with no founded or well-founded order (e.g. `a :- know b. b :- know a.`), and unbounded tuple-key forms (cons `[H|T]`, predicate references, aggregates), each rejected at compile time with a typed diagnostic rather than silently approximated.

7.2 World-View Semantics

A world view is a non-empty set of stable models. `know p` holds when `p` appears in every model of the world view; `possible p` holds when it appears in at least one; `not know` and `not possible` are their negations. Two semantics are selectable through `#pragma epistemic_mode`. The default, FAEEL (Founded Autoepistemic Equilibrium Logic), rejects circular epistemic support unless it is independently founded, so a self-supporting rule such as `p :- possible p.` fails closed. A Gelfond-1991-style compatibility mode, `g91`, instead admits such self-support for legacy epistemic specifications. Non-epistemic programs evaluate identically under both modes.

7.3 GPU-Native Execution

Accepted epistemic programs lower from EIR to a GPU-executable plan rather than to ordinary RIR. Execution follows a Generate–Propagate–Test schedule whose three phases run as CUDA kernels: *generate* enumerates bounded candidate assumptions as device bitsets; *propagate* prunes candidates that immediately contradict known or rejected facts; and *test* validates the survivors against the program’s world views, checking stable-model tuple membership on the GPU for predicates of each arity. The reduced ordinary part of each rule reuses the existing optimizer, WCOJ planner, and helper-splitting machinery (Section 4), and *epistemic splitting* decomposes a program into independent components that are solved separately and recombined. The accepted hot path performs no CPU fallback—candidate enumeration, membership tests, and world-view validation all execute on device—and a transfer budget tracks that no data-plane host transfers occur within it.

Recursive epistemic programs execute through the same substrate. A modal atom that ranges over an invariant relation reduces to its gated relation and runs through the ordinary GPU recursive fixpoint; a modal over an epistemically *determined* head is resolved by stratified execution (the theorem $\text{know } R \equiv R$ applied at the relation-store boundary); positive recursion computes the FAEEL founded least fixpoint, or the G91 self-support fixpoint under **g91**; and a negated-modal cycle with no stratified order is resolved by a GPU-backed well-founded-semantics (WFS) alternating fixpoint rather than the host probabilistic WFS solver. That WFS path is GPU-backed but still host-orchestrated—it is not yet a device-resident, no-host-interaction WFS, which remains future work (Section 11).

7.4 Solver and Probabilistic Integration

Epistemic candidate validation reuses the on-GPU CDCL solver already built for knowledge-compilation verification (Section 5): the same device-resident CNF/CDCL substrate handles satisfiability checks, bounded MaxSAT, and assumption-based queries, avoiding a separate epistemic SAT engine. Accepted world-view evidence likewise feeds the existing GPU-native exact/provenance path, so a query can combine epistemic and probabilistic structure without leaving the device. In the v0.9.x line the GPU-native epistemic execution path is the mature surface; the broader solver portfolio and the end-to-end probabilistic–epistemic evidence path are still being completed (Section 11).

8 Interoperability

A GPU-native logic engine is only as useful as its ability to exchange data with the frameworks researchers already use. *xlog* exposes GPU-resident relations, probabilities, and gradients through four standard interfaces—DLPack, Arrow, Python bindings, and PyTorch autograd—so that results flow into downstream tools without device-to-host copies or format conversion.

8.1 DLPack: Zero-Copy GPU Tensor Sharing

xlog implements the DLPack [8] protocol for bidirectional zero-copy exchange of GPU buffers with any framework that speaks it—PyTorch, JAX, CuPy, cuDF. On export, each relation column becomes a managed DLPack tensor whose underlying GPU memory is reference-counted and freed

only once every consumer releases its handle. On import, xlog validates device affinity, dtype, contiguity, and alignment—optionally against an expected schema, rejecting mismatches before evaluation—and assembles the tensors into a device buffer with no host copy. All eight scalar types have well-defined DLPack dtype mappings. In Python the protocol surfaces as a PyCapsule, so a consumer calls `torch.from_dlpack(capsule)` to obtain a live CUDA tensor backed by the very memory xlog wrote.

8.2 Arrow IPC: Columnar Export with Dictionary Encoding

For columnar analytics tools—Pandas, Polars, DuckDB—xlog exports relations through the Arrow [9] C Device Data Interface, pairing Arrow array and schema pointers with a CUDA device type and ordinal for zero-copy round-tripping. Symbol-typed columns get special handling: interned symbol IDs export as an Arrow dictionary array over a string dictionary, so a Polars or DuckDB consumer reads symbolic columns as native string dictionaries with no separate lookup, while xlog keeps the compact integer representation internally.

8.3 Python Bindings

The `pyxlog` crate provides Python bindings built with PyO3 [11], exposing compilation entry points for deterministic, probabilistic, and ILP programs, each returning a compiled object with evaluation, training, and result-extraction methods; type stubs give full IDE auto-completion and static checking. Tensor-valued results are returned as DLPack capsules, keeping data on the GPU and letting the caller choose how to materialize it. Long-running GPU operations release the Python GIL before kernel dispatch—circuit evaluation, ILP fixpoint iteration, and the neural backward pass all run without holding it, so Python data-loader threads and asyncio tasks proceed concurrently with GPU computation.

8.4 PyTorch Autograd Integration

Neural-symbolic training requires gradients to flow from logical inference back into neural network parameters. xlog achieves this by constructing loss tensors that participate directly in PyTorch’s autograd graph: it evaluates the compiled circuit on the GPU—with the GIL released—then exports the scalar NLL loss as a DLPack capsule and imports it via `torch.from_dlpack`. Because the network’s forward pass produced grad-enabled tensors that fed into the circuit, the loss tensor stays connected to the autograd graph. A batched path stacks per-network inputs into one forward pass, evaluates multiple queries against the shared circuit, accumulates per-query losses, and runs a single `torch.autograd.backward`. The circuit thus behaves as a differentiable function inside PyTorch—gradients flow from the logic layer through DLPack tensors back to `nn.Module` parameters, with no custom autograd `Function` subclass required.

9 Evaluation

This section reports a mix of artifact-backed measurements and benchmark acceptance targets for xlog’s deterministic, probabilistic, and neural-symbolic subsystems. Every quantitative claim is tied either to a checked-in artifact or to the benchmark contract in `docs/BENCHMARKS.md`; the two are

distinguished explicitly below. These are single-system measurements on development hardware, not controlled head-to-head comparisons against other engines.

9.1 Methodology

Hardware. All measurements were collected on an NVIDIA RTX PRO 3000 Blackwell Generation Laptop GPU (12 GB VRAM, SM120, compute capability 12.0, driver 591.59) (`docs/BENCHMARKS.md`).

Build configuration. Rust crates are compiled in `-release` mode. CUDA kernel modules are staged cubin-first with portable PTX fallback and explicit JIT warmup. Python components use `pyxlog` with CUDA-enabled PyTorch.

Statistical protocol. The benchmark harness uses Criterion.rs with the settings in Table 3 (`docs/BENCHMARKS.md`).

Table 3: Criterion.rs benchmark settings.

Setting	Value
Sample size	10–100 runs (benchmark-dependent)
Warm-up	3 iterations (PTX JIT, memory pool init)
Significance level	0.10
Noise threshold	0.05 (ignore <5% variance)
Seeding	Deterministic LCG (no system entropy)

Artifact-backed values are quoted directly from checked-in JSON/Markdown evidence. Tables sourced from `docs/BENCHMARKS.md` are benchmark targets for the development machine rather than archived run logs. GPU measurements include explicit warm-up to amortize kernel-module loading and CUDA context establishment.

9.2 Absolute Performance

9.2.1 Deterministic Subsystem

For transitive closure, hash join, and aggregation in isolation the repository publishes target throughput envelopes (Section A) rather than archived result tables, so we make no measured absolute-throughput claim for these operations here. The measured deterministic result reported in this paper is the worst-case-optimal join ablation below.

9.2.2 Worst-Case Optimal Joins

The WCOJ subsystem (Section 4) is measured against the binary-join baseline on four paper-scale fixtures (call-graph edges, Andersen points-to [18], `ddisasm`-style disassembly [19], and a neural-symbolic mining analog), each with 4,194,304 input rows. The first three are standard Datalog program-analysis shapes; the fourth is a recursive neural-symbolic rule-mining workload. We measure at this scale deliberately: WCOJ speedups depend strongly on input size and skew, and synthetic micro-benchmarks systematically overstate them, so multi-rule recursive workloads at millions of tuples are the honest setting for a reproducible claim. Table 4 reports the artifact-backed speedups;

both paths are required to produce identical row sets before timing, and every measurement is a median over 10 runs with coefficient of variation ≤ 0.05 .

Table 4: WCOJ vs. binary-join speedup on paper-scale fixtures (4.19M input rows each). Source: [docs/evidence/2026-05-18-g39-w39-dts-dlm-analog/](#).

Fixture	WCOJ speedup
<code>call_graph_edge</code>	29.62×
<code>andersen</code>	26.96×
<code>ddisasm</code>	28.79×
<code>nesy_analog</code>	26.60×
Geometric mean	27.96×

These are single-system measurements of xlog’s own WCOJ path against its own binary-join path on skewed inputs; on uniform or empty inputs the adaptive classifier falls back to the binary-join chain. They are an ablation, not a comparison against another engine—this paper makes no head-to-head claim against any external system (Section 11). Peak observed allocation was roughly 2.29 GB, within the device’s 12 GB budget.

9.2.3 Runtime Optimization Layer

The v0.8.6 runtime optimizations (Section 4) are measured on dedicated repeated-session fixtures. The persistent hash-index manager records a 3.21× speedup on a build-heavy repeated-session semi-join (cached median 0.079s vs. uncached 0.255s) against a 1.5× target, with zero tracked device-to-host or host-to-device transfers ([docs/evidence/2026-05-19-v086-persistent-hash-index/](#)). The profile-gated shared-memory chain scorer records a 5.58× speedup on a chain-hot fixture, with a bounded −1.3% regression on the small-input control ([docs/evidence/2026-05-19-v086-chain-smem/](#)). Common-subexpression elimination and adaptive re-optimization are validated for output equivalence and duplicate-work reduction with no added data-plane transfers, but are reported here as correctness-gated optimizations rather than headline speedups.

9.2.4 Probabilistic Subsystem

The probabilistic engine supports exact inference via Decision-DNNF knowledge compilation and GPU-accelerated Monte Carlo sampling. The repository publishes target latency and throughput envelopes for both (Section A) but does not yet ship archived absolute-timing reports with stable per-phase numbers, so we make no measured absolute-latency claim for the engine in isolation. The measured probabilistic result in this paper is the circuit-caching ablation of the next subsection. The MC engine’s hot-loop optimizations—targeted reset planning, bulk GPU sampling, and pointer-buffer reuse—are present in the implementation but unquantified here.

9.2.5 Neural-Symbolic Subsystem

Neural-symbolic training is measured on the `01_minimal` benchmark (MNIST single-digit addition, 512 training images, 5 epochs, batch size 64). The timing summary below uses the latest checked-in

Table 5: Training performance summary.

Metric	Value	Source
PTX JIT cold start	0.02s	Cubin-first module loading target from <code>docs/BENCHMARKS.md</code>
First epoch (cold)	71.654s	Track B comparison, mean over 3 seeds
Steady-state epoch (warm)	0.248s	Track B comparison, mean epochs 2–5
Total training (5 epochs)	72.645s	Track B comparison, mean over 3 seeds
Per-query cost	0.97ms	Track B comparison, mean forward+backward query cost
Circuit cache speedup	2.74×	Cache ablation artifact, 95% CI: [2.29, 3.18]

post-fix comparison artifact on current `main`; the cache-ablation study remains a separate 3-epoch measurement of the compilation-cache effect.

The cache ablation study compared circuit-cached vs. uncached training over 3 seeds and 3 epochs each. Mean training time dropped from 242.9s (uncached) to 88.9s (cached), yielding the 2.74× speedup reported in Section 5 (`examples/neural/results/evidence/cache_ablation_20260218.json`).

DeepProbLog baseline. The checked-in DeepProbLog comparison artifact is partial and protocol-mismatched, so it should be read as a qualitative baseline note rather than a clean head-to-head measurement. In `examples/neural/results/track_a/20260216T145409Z_track_a_dev/comparisons/mnist_vs_deepproblog.json`, the strict full-train DeepProbLog configuration (2 epochs, batch size 2048) reached held-out addition accuracy of 23.24% after 2:31:07, while a smaller limited run was manually aborted after 60 minutes. The same artifact reports XLOG training loss rather than held-out accuracy and marks one XLOG seed incomplete, so direct quantitative comparison remains premature.

GA-based ILP. The genetic algorithm reliability gate (`test_ga_reliability_50`) achieved a 3.25× wall-clock reduction from 1447s to 445.30s by reducing redundant attempts, with identical statistical quality: 200/200 success rate and Clopper–Pearson lower-95 CI = 0.981725 (`docs/evidence/2026-03-05-ga-runtime-closure/sweep-results.md`).

9.3 Qualitative Comparison

Table 6 positions xlog relative to existing systems across key capability dimensions. Each system addresses a subset of the design space; xlog’s contribution is unifying these capabilities within a single GPU-resident runtime.

Published GPU Datalog results. GPUlog [3] reports up to 45× speedup over CPU baselines using HISA (Hash-Indexed Sorted Array) indexing with lock-free deduplication. VFLog [4] demonstrates up to 200× gains over CPU column engines with a column-oriented GPU Datalog runtime. Both systems are restricted to deterministic evaluation; neither supports probabilistic inference, knowledge compilation, or neural network integration. xlog draws on the same insight—that Datalog’s relational algebra maps naturally to GPU parallelism—but extends the execution model through the probabilistic and neural-symbolic layers described in Sections 5 and 6.

Honest framing. The comparison above is qualitative. xlog has not been benchmarked head-to-head against GPUlog or VFLog on identical workloads. The deterministic throughput targets (Section A) represent xlog’s own benchmark contract, not comparative claims, and the WCOJ and circuit-caching numbers are ablations against xlog’s own baselines. The DeepProbLog comparison

Table 6: Capability comparison across probabilistic and GPU logic systems.

Dimension	DeepProbLog	ProbLog2	GPUlog	xlog
Symbolic execution	CPU (SWI-Prolog)	CPU (tabling)	GPU (HISA)	GPU (semi-naive fixpoint)
Probabilistic inference	Exact (PySDD, CPU)	Exact + approx (CPU)	None	Exact (d-DNNF, GPU) + MC (GPU)
Knowledge compilation	PySDD (CPU)	d-DNNF via c2d/D4 (CPU)	None	D4 compile + GPU circuit eval
Neural integration	PyTorch + Prolog bridge	None	None	Zero-copy DLPack, fused backward
Zero-copy ML interop	No (CPU-GPU transfers)	No	No	Yes (DLPack, device-resident tensors)
Differentiable ILP	Via TP-operator (CPU)	No	No	GPU-resident credit/loss path (beta)

is limited to the `01_minimal` task with differing dataset sizes, metrics, and partially completed archived runs, as noted above.

9.4 CUDA Certification

The CUDA kernel certification suite (`crates/xlog-cuda-tests/`) provides systematic coverage of all GPU kernel operations. On current `main`, the statistics in Table 7 hold.

Table 7: CUDA certification suite statistics (current `main`, May 2026).

Metric	Value
Total tests	207
Categories	33 (C01–C25 infrastructure + G01–G08 GPU-specific)
Latest full certification	100% (207/207)
Kernel modules covered	22

The 25 infrastructure categories (C01–C25) cover toolchain validation, launch configuration, pointer bounds, memory spaces (global, shared, local), synchronization, warp-level operations, control flow divergence, atomics, floating-point semantics, integer operations, determinism, async pipelines, caching, host-device transfer, multi-stream execution, multi-GPU, hardware stress, algorithmic correctness, blind spots, edge matrices, and float filter semantics (`docs/certification/2026-01-22-v0.4.0-alpha-certification-report.md`). The 8 GPU-specific categories (G01–G08) validate circuit forward evaluation, circuit backward (gradient) kernels, weight injection, transfer efficiency, circuit caching, PTX robustness, SAT/CDCL solving, and device-count invariants (`docs/architecture/cuda-certification.md`). The 22 covered runtime kernel modules correspond to the staged runtime CUDA kernel set plus the build-time ILP-family modules documented in the certification architecture notes. Because the public repository does not rely on hosted GPUs, the full certification suite is exercised through GPU-backed local release validation (`scripts/validate_release_gpu.sh -mode release`) and the manual self-hosted `cuda-ci` workflow; the always-on

push/PR CI remains non-GPU and covers formatting, shell/workflow linting, package metadata, and no-GPU build checks.

10 Related Work

xlog draws on and extends five lines of research: the tradition of logic programming languages, GPU-accelerated Datalog evaluation, probabilistic logic programming, GPU SAT solving, and differentiable inductive logic programming. We survey each in turn, positioning xlog’s contributions relative to the state of the art.

10.1 Logic Programming Languages

xlog sits in a long line of logic programming languages but is the first to target GPU execution of the full language surface. Prolog [20] is the canonical universal logic language, offering unification and extensive metaprogramming, running on a CPU tree-walking interpreter. Mercury [21] introduced strong typing and mode analysis to logic programming, compiling to optimized native code; xlog inherits the typed-predicate idea but targets device kernels rather than CPU code. Soufflé [22] is a typed Datalog dialect that compiles to efficient C++ and is widely used for program analysis; it is CPU-resident. Answer-set programming with clingo [23] provides nonmonotonic reasoning via answer-set semantics on CPU grounders and solvers; xlog is not an ASP system, but its stratified semantics and probabilistic extensions intersect ASP applications. Probabilistic logic languages such as ProbLog [15] and DeepProbLog [1] express weighted facts and neural predicates as first-class syntactic constructs, but their reasoners run on the CPU while neural components run on the GPU. xlog unifies all of these—typed logic, Datalog evaluation, probabilistic facts, neural predicates—as one language, compiled and executed entirely on the GPU.

10.2 GPU Datalog

Prior GPU-native Datalog systems—gpulog [3] and vflog [4]—target the Datalog evaluation backend that xlog also provides (Section 4), but do not surface a language of xlog’s breadth (typed predicates, UDFs, modules, probabilistic and neural extensions). GPUlog ported bottom-up fixed-point computation to CUDA and reported approximately $45\times$ speedups over equivalent CPU Datalog engines on transitive-closure benchmarks. VFLog advances this line with a vertically-fused, column-oriented kernel design that avoids materializing intermediate relations, achieving up to $200\times$ speedups over CPU column engines on recursive graph queries. More recently, mnmgDatalog [24] extends GPU Datalog to multi-node, multi-GPU clusters, distributing semi-naive evaluation across devices via partitioned fact stores and all-to-all delta exchanges. All three systems target deterministic Datalog: they implement relational algebra operators (join, project, deduplicate) on the GPU and accelerate fixed-point iteration, but none supports probabilistic semantics, weighted model counting, or gradient computation over derived facts. xlog shares VFLog’s commitment to columnar storage and kernel fusion for relational operators, but extends the GPU-resident execution model to probabilistic and differentiable reasoning, maintaining a single address space from semi-naive evaluation through circuit-based inference to neural-symbolic gradient propagation. Where these prior GPU engines rely on binary joins, xlog additionally promotes eligible cyclic and multiway

rules to a worst-case-optimal join path [12–14] (Sections 4 and 9), avoiding the intermediate-relation blowup that binary-join plans incur on skewed cyclic queries.

10.3 Probabilistic Logic Programming

ProbLog2 [15] established the modern pipeline for exact probabilistic inference in logic programs: ground the relevant program, encode provenance as a propositional formula, compile to a tractable target (d-DNNF or SDD), and evaluate weighted model counts over the compiled circuit. DeepProbLog [1] extends ProbLog2 by replacing selected probabilistic facts with the outputs of neural networks, enabling end-to-end gradient-based training of neural predicates within a logical framework. NeurASP [2] takes a related approach, integrating neural network outputs as probability annotations on answer-set programs and using stable-model semantics for inference. These systems demonstrate that combining neural perception with symbolic reasoning yields strong performance on tasks requiring both pattern recognition and logical structure. However, all three perform symbolic inference on the CPU: ProbLog2 invokes the D4 or c2d [25] compiler as a host-side subprocess, DeepProbLog inherits this pipeline and adds a Python-level training loop that shuttles gradients between CPU logic and GPU tensors, and NeurASP relies on CPU-bound answer-set solving. The checked-in DeepProbLog baseline artifact in this repository is qualitative rather than a clean head-to-head benchmark: it records a 2:31:07 full-train run at 23.24% held-out addition accuracy and a smaller aborted run, while also noting protocol mismatches versus XLOG’s archived measurements (`examples/neural/results/track_a/20260216T145409Z_track_a_dev/comparisons/mnist_vs_deepproblog.json`). xlog eliminates the host/device bottleneck by compiling d-DNNF circuits into GPU-resident XGCF format with compile-once/evaluate-many semantics, executing both forward weighted model counting and backward gradient propagation as level-parallel CUDA kernels without host transfers during training iterations.

10.4 GPU SAT

ParaFROST [5] is a GPU-accelerated CDCL SAT solver that parallelizes clause database simplification—bounded variable elimination, subsumption, and self-subsuming resolution—on the GPU while retaining sequential conflict-driven search on the CPU. It is a complete solver with proof-generation capabilities, competitive with MiniSat and CaDiCaL on industrial benchmarks. FastFourierSAT [26] takes a different approach, reformulating SAT as continuous optimization over Fourier coefficients and applying local search on the GPU; it is an incomplete solver suited to satisfiable instances but unable to certify unsatisfiability. Both systems are standalone solvers designed for the SAT competition setting, not embedded within a logic programming pipeline. xlog’s GPU CDCL module (Section 5, Stage 4) shares ParaFROST’s strategy of offloading clause manipulation to device kernels, but integrates the solver as a verification component within the knowledge compilation pipeline—checking circuit-formula equivalence via two UNSAT proofs, computing minimal unsatisfiable cores, and providing proof certificates that ensure the correctness of compiled circuits.

10.5 Differentiable ILP

Differentiable inductive logic programming (dILP) [6] demonstrates that first-order rule induction can be cast as a differentiable optimization problem: candidate rules are scored by soft forward-

chaining over all possible rule instantiations, and gradient descent adjusts rule weights to maximize coverage of positive examples. The approach is elegant but computationally demanding: forward-chaining materializes a dense tensor over the full rule space at each step, scaling cubically with the number of constants in the domain. Subsequent work has reduced this cost through pruning and approximate materialization, but implementations remain CPU-based and struggle beyond a few hundred constants. xlog’s dILP subsystem (Section 6) replaces dense rule materialization with sparse GPU mask computation: candidate rules are represented as sparse bitmasks over the ground-atom space, credit assignment is performed entirely on-device via scatter-gather operations, and a six-gate promotion pipeline controls rule acceptance. This sparse, GPU-resident formulation avoids the cubic blowup of dense materialization while preserving differentiability, enabling rule induction over domains with thousands of constants.

10.6 Positioning

Each system surveyed above excels in its niche. GPUlog and VFLog demonstrate that relational fixed-point computation maps naturally to GPU architectures. ProbLog2 and DeepProbLog prove that knowledge compilation provides a principled bridge between logic and learning. ParaFROST shows that clause-level parallelism yields practical SAT speedups. dILP establishes that rule induction is amenable to gradient-based optimization. xlog unifies these capabilities on a single GPU-resident platform with zero-copy ML interop—not claiming superiority on any single axis, but providing the integrated stack. By keeping fact stores, compiled circuits, solver state, and gradient buffers in device memory throughout the reasoning pipeline, xlog eliminates the host-device transfer overhead that fragments existing multi-system approaches, offering practitioners a single engine where deterministic evaluation, probabilistic inference, SAT verification, and differentiable learning compose without leaving the GPU.

11 Limitations and Future Work

11.1 Current Limitations

NVIDIA GPU required. xlog targets CUDA exclusively. The entire execution stack—from hash-join kernels and radix sort to Monte Carlo sampling and knowledge compilation—is written in CUDA C. There is no OpenCL, Metal, or ROCm backend. This is a direct consequence of the GPU-native design philosophy: rather than abstracting over heterogeneous devices behind a portability layer, xlog commits to a single GPU programming model to exploit CUDA-specific features such as cooperative groups, warp-level primitives, and unified virtual addressing. These primitives are load-bearing, not incidental: the persistent-index manager, the recorded-launch discipline, and the in-kernel hash-consing of the knowledge-compilation pipeline all depend on them, and a lowest-common-denominator portability layer—forcing extra synchronization points and disallowing warp operations—would erode the zero-transfer performance model that is the system’s reason for existing. Portability is thus traded away deliberately. Users without an NVIDIA GPU cannot run xlog.

All data must fit in GPU memory. Relations, indices, circuit buffers, and Monte Carlo sample arrays are allocated entirely on-device. There is no out-of-core spilling mechanism; a program whose working set exceeds GPU VRAM will fail with a `RESOURCE_EXHAUSTED` error rather than

degrade silently. For most Datalog workloads this is not a bottleneck—a 24 GB GPU holds billions of 8-byte tuples—but large-scale knowledge graphs or high-sample-count MC inference can hit the ceiling.

Python batch-query helpers are typed but not float-complete. The Python helper APIs—`batch_fact_membership`, `batch_fact_membership_device`, `batch_tagged_credit`, and `batch_tagged_credit_device`—now perform schema-aware uploads for `u32`, `u64`, `i32`, `i64`, and `bool` columns, so wide integer IDs no longer truncate silently. However, these helper paths still reject floating-point batch uploads (`f32/f64`) at the binding layer. The underlying engine and DLPack relation-import paths remain typed; this is a convenience-API gap rather than a core execution limit.

Differentiable ILP is beta. The dILP subsystem—structure learning via gradient-based rule search—shipped in v0.4.0 and received significant reliability and performance work from v0.5.0 through v0.9.2. It remains beta: the search space is restricted to definite programs, convergence is sensitive to learning-rate schedules, and the API surface may change in minor releases.

Epistemic and solver maturity. The epistemic runtime (Section 7) executes accepted programs GPU-natively in the v0.9.2 release, including finite nested modal chains, epistemic integrity constraints, multi-predicate epistemic coupling, same-name multi-arity modal predicates, and recursive epistemic execution. What remains immature is the broader solver surface rather than the epistemic core: the MaxSAT/portfolio solver services and the end-to-end probabilistic–epistemic evidence path ship as a bounded surface still being broadened, and a dedicated Solver IR (SIR) remains planned (ROADMAP.md). Genuine fail-closed boundaries persist by design—raw RIR lowering of modal literals, genuinely cyclic modal coupling with no founded or well-founded order, and unbounded tuple-key forms—each rejected with a typed diagnostic (Section 7).

No formal head-to-head benchmarks. xlog’s probabilistic inference pipeline occupies similar territory to DeepProbLog, but no controlled comparison on identical programs and datasets exists. The baseline experiments in Section 9 use different dataset sizes and evaluation protocols. Claims of speedup relative to CPU-based probabilistic logic systems remain informal until such benchmarks are published.

11.2 Future Work

The following directions remain outside the current released contract.

Out-of-core execution. For programs whose materialized relations exceed GPU memory, a streaming execution mode would partition relations across host and device, paging tiles through GPU memory in fixpoint-iteration order. This would remove the single-GPU-memory ceiling described above, at the cost of added host-device transfer overhead. The v0.8.6 persistent hash-index manager reduces repeated-session build cost but does not spill to host; out-of-core execution is targeted for v0.10.0 alongside multi-GPU partitioning (ROADMAP.md).

Epistemic residency expansion. The v0.9.x line implements an accepted EIR/GPU epistemic execution surface with FAEEL default semantics, G91 compatibility mode, finite nested modal-chain normalization, split/joint planning, tuple-key membership, and GPU-backed runtime paths. The remaining future work is stronger residency and coverage: device-resident/no-host-interaction WFS, broader non-finite tuple-key forms, and semantic forms outside the finite accepted release surface.

Multi-GPU partitioned evaluation. Inspired by mnmgDatalog’s radix-hash partitioning

and GPU-aware all-to-all shuffle, a multi-GPU backend would distribute relations across devices and coordinate join and union operations via NVLink or PCIe. Partitioning kernels are described in the roadmap but unimplemented; this is targeted for v0.10.0 (`ROADMAP.md`).

A Performance Contract

The tables in this appendix are *benchmark targets*, not measured results. They define the throughput envelopes the deterministic and probabilistic subsystems are expected to meet on the development GPU (Section 9) and are published as a regression contract in `docs/BENCHMARKS.md`. They are kept separate from Section 9, which reports only artifact-backed measurements, precisely so the two are not conflated. Actual throughput depends on graph structure, key distribution, and output cardinality.

Table 8: Deterministic throughput targets.

Operation	Configuration	Target
Transitive closure	100K random edges (sparse)	>1M rows/sec
Transitive closure	1M random edges (medium)	>5M rows/sec
Transitive closure	$K_{500,500}$ bipartite (dense)	>10M rows/sec
Hash join	$100K \times 100K$	>50M rows/sec
Hash join	$1M \times 100K$	>100M rows/sec
Hash join	High selectivity	>20M rows/sec
Aggregation	1M rows, 100K groups	Tracked; no fixed threshold

Table 9: Exact inference targets.

Configuration	Target	Notes
20-variable probabilistic path	<100 ms	Small d-DNNF circuit
50-variable Bayesian network	<500 ms	Medium circuit complexity
With gradient computation	<2× base time	Backward pass overhead

Table 10: Monte Carlo inference targets.

Configuration	Target	Notes
100K samples, 100 variables	>10M worlds/sec	Throughput mode
10K samples, 500 variables	>5M worlds/sec	Complexity mode

References

- [1] Robin Manhaeve et al. “DeepProbLog: Neural Probabilistic Logic Programming”. In: *Advances in Neural Information Processing Systems (NeurIPS)*. 2018.
- [2] Zhun Yang, Adam Ishay, and Joohyung Lee. “NeurASP: Embracing Neural Networks into Answer Set Programming”. In: *Proceedings of the Twenty-Ninth International Joint Conference on Artificial Intelligence (IJCAI)*. 2020.

- [3] Yihao Sun et al. “Optimizing Datalog for the GPU”. In: *Proceedings of the 30th ACM International Conference on Architectural Support for Programming Languages and Operating Systems (ASPLOS)*. 2025. DOI: [10.1145/3669940.3707274](https://doi.org/10.1145/3669940.3707274).
- [4] Yihao Sun et al. “Column-Oriented Datalog on the GPU”. In: *Proceedings of the AAAI Conference on Artificial Intelligence*. Vol. 39. 14. 2025, pp. 15177–15185. DOI: [10.1609/aaai.v39i14.33665](https://doi.org/10.1609/aaai.v39i14.33665).
- [5] Muhammad Osama, Anton Wijs, and Armin Biere. “SAT Solving with GPU Accelerated Inprocessing”. In: *Tools and Algorithms for the Construction and Analysis of Systems (TACAS)*. Vol. 12651. Lecture Notes in Computer Science. Springer, 2021, pp. 133–151. DOI: [10.1007/978-3-030-72016-2_8](https://doi.org/10.1007/978-3-030-72016-2_8).
- [6] Richard Evans and Edward Grefenstette. “Learning Explanatory Rules from Noisy Data”. In: *Journal of Artificial Intelligence Research* 61 (2018), pp. 1–64.
- [7] Jean-Marie Lagniez and Pierre Marquis. “An Improved Decision-DNNF Compiler”. In: *Proceedings of the Twenty-Sixth International Joint Conference on Artificial Intelligence (IJCAI)*. 2017.
- [8] *DLPack: Open In Memory Tensor Structure*. <https://github.com/dmlc/dlpack>. Accessed 2026.
- [9] *Apache Arrow: Cross-Language Development Platform for In-Memory Data*. <https://arrow.apache.org>. Accessed 2026.
- [10] *cudarc: Safe Rust Wrapper around CUDA*. <https://github.com/coreylowman/cudarc>. Accessed 2026.
- [11] *PyO3: Rust Bindings for Python*. <https://pyo3.rs>. Accessed 2026.
- [12] Hung Q. Ngo et al. “Worst-case Optimal Join Algorithms”. In: *Journal of the ACM* 65.3 (2018), 16:1–16:40.
- [13] Todd L. Veldhuizen. “Leapfrog Triejoin: A Simple, Worst-Case Optimal Join Algorithm”. In: *International Conference on Database Theory (ICDT)*. 2014, pp. 96–106.
- [14] Yisu Remy Wang, Max Willsey, and Dan Suciu. “Free Join: Unifying Worst-Case Optimal and Traditional Joins”. In: *Proceedings of the ACM on Management of Data (SIGMOD)* 1.2 (2023), 150:1–150:23.
- [15] Daan Fierens et al. “Inference and Learning in Probabilistic Logic Programs using Weighted Boolean Formulas”. In: *Theory and Practice of Logic Programming* 15.3 (2015), pp. 358–401.
- [16] Grigori S. Tseitin. “On the Complexity of Derivation in Propositional Calculus”. In: *Studies in Constructive Mathematics and Mathematical Logic* (1968).
- [17] João P. Marques-Silva and Karem A. Sakallah. “GRASP—A New Search Algorithm for Satisfiability”. In: *Proceedings of the IEEE/ACM International Conference on Computer-Aided Design (ICCAD)*. 1996.
- [18] Lars Ole Andersen. “Program Analysis and Specialization for the C Programming Language”. PhD thesis. University of Copenhagen, 1994.
- [19] Antonio Flores-Montoya and Eric Schulte. “Datalog Disassembly”. In: *USENIX Security Symposium*. 2020, pp. 1075–1092.

- [20] Jan Wielemaker et al. “SWI-Prolog”. In: *Theory and Practice of Logic Programming* 12.1-2 (2012), pp. 67–96.
- [21] Zoltan Somogyi, Fergus Henderson, and Thomas Conway. “The Execution Algorithm of Mercury, an Efficient Purely Declarative Logic Programming Language”. In: *Journal of Logic Programming* 29.1-3 (1996), pp. 17–64.
- [22] Herbert Jordan, Bernhard Scholz, and Pavle Subotić. “Soufflé: On Synthesis of Program Analyzers”. In: *Computer Aided Verification (CAV)*. Springer, 2016, pp. 422–430.
- [23] Martin Gebser et al. “Multi-shot ASP Solving with clingo”. In: *Theory and Practice of Logic Programming* 19.1 (2019), pp. 27–82.
- [24] Ahmedur Rahman Shovon et al. “Multi-Node Multi-GPU Datalog”. In: *Proceedings of the 39th ACM International Conference on Supercomputing (ICS)*. 2025. DOI: [10.1145/3721145.3730431](https://doi.org/10.1145/3721145.3730431).
- [25] Adnan Darwiche. “New Advances in Compiling CNF into Decomposable Negation Normal Form”. In: *Proceedings of the 16th European Conference on Artificial Intelligence (ECAI)*. 2004.
- [26] Yunuo Cen, Zhiwei Zhang, and Xuanyao Fong. *Massively Parallel Continuous Local Search for Hybrid SAT Solving on GPUs*. arXiv:2308.15020. 2023.